M.Tech Program

Advanced Industry Integrated Programs

Jointly offered by University and LTIMindTree

Generative AI Fundamentals: Apply Generative Adversarial Networks (GANs)

Knowledge partner

LTIMindtree

Implementation partner

L&T EduTech

Modules to cover….

1. Build Basic Generative Adversarial Networks
2. Case Study 1
3. Build Better Generative Adversarial Networks
4. Case Study 2
5. **Apply Generative Adversarial Networks**
6. Case Study 3

Apply Generative Adversarial Networks

Apply Generative Adversarial Networks

GANs for Data Augmentation and Privacy Preservation

Why GANs for Data Augmentation?

- GANs can generate large amounts of synthetic data, especially useful when real-world data is scarce.
- Used to enhance datasets for training machine learning models, improving accuracy and robustness.

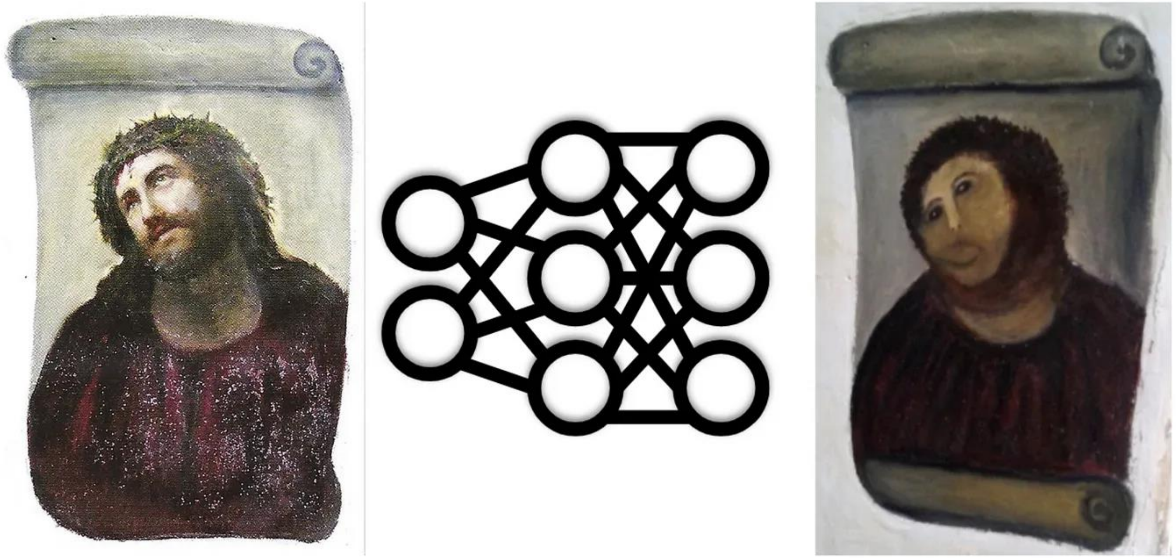

Image Source: https://medium.com/abacus-ai/gans-for-data-augmentation-21a69de6c60b

Apply Generative Adversarial Networks

GANs for Data Augmentation and Privacy Preservation

- **Create Synthetic Data to Augment Existing Datasets**
  - ✓ **Example:** Generating additional images of objects in different angles or lighting conditions.
  - ✓ Boosts model performance by providing a more varied dataset.

- **Privacy-Preserving GANs**
  - ✓ GANs can generate data that mimics real-world data without revealing sensitive information.
  - ✓ Prevents privacy leaks, especially in domains where data sensitivity is a concern.

Image Source: https://medium.com/abacus-ai/gans-for-data-augmentation-21a69de6c60b

Apply Generative Adversarial Networks

GANs for Data Augmentation and Privacy Preservation

- **Use Cases for Ensuring Privacy through Synthetic Data Generation**
  - ✓ Protecting patient information in healthcare by generating synthetic medical records.
  - ✓ Creating financial datasets that protect individual user data while preserving overall trends.
- **Examples**:
  - ✓ **Healthcare**: Generating synthetic patient data for medical research without compromising patient privacy.
  - ✓ **Finance**: Creating financial transaction data that maintains statistical properties without exposing sensitive user information.

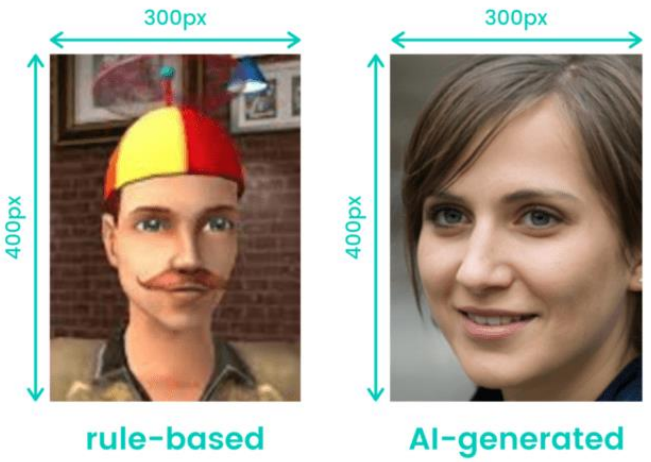

Image Source: https://mostly.ai/blog/define-synthetic-data

Applications of GANs for Data Augmentation, Privacy, and Anonymity

Apply Generative Adversarial Networks

Applications of GANs for Data Augmentation, Privacy, and Anonymity

- **Data Augmentation with GANs**

  ✓ GANs have become a powerful tool for data augmentation, especially in scenarios where acquiring large, diverse datasets is difficult. By generating synthetic data, GANs can expand the available training data to improve the performance of machine learning models.

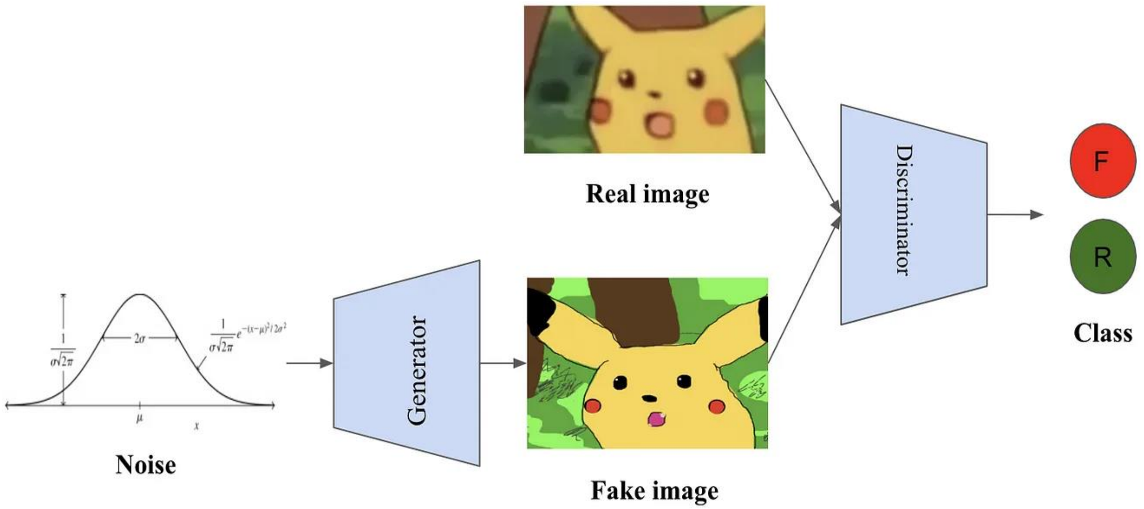

Image Source: https://medium.com/abacus-ai/gans-for-data-augmentation-21a69de6c60b

Apply Generative Adversarial Networks

Applications of GANs for Data Augmentation, Privacy, and Anonymity

- **Example Use Cases**:
  - ✓ **Medical Imaging**: GANs are used to create realistic medical scans (e.g., MRI, CT) for training AI models, improving detection of anomalies like tumors.
  - ✓ **Facial Recognition**: GANs can generate a variety of facial images with different lighting, expressions, or angles, helping improve the accuracy of facial recognition systems.
  - ✓ **Autonomous Vehicles**: In autonomous driving, GANs generate synthetic driving scenarios (like different weather conditions, lighting, or road settings) to train vehicle systems for better performance in real-world situations.

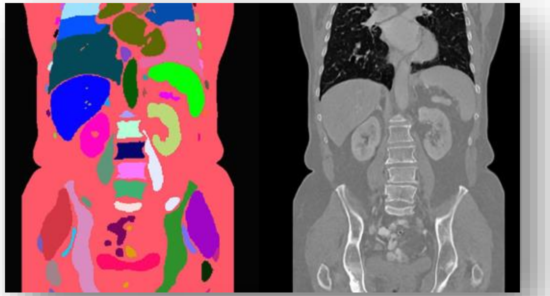

Image Source: https://developer.nvidia.com/blog/addressing-medical-imaging-limitations-with-synthetic-data-generation/
Image Source: https://www.thesun.co.uk/tech/7544615/facial-recognition-cameras-policing-britain/
Image Source: https://www.wired.com/story/google-self-driving-car-stop-development/

Apply Generative Adversarial Networks

Applications of GANs for Data Augmentation, Privacy, and Anonymity

- **Role in Privacy Preservation:**
  - ✓ GANs can help in privacy-preserving scenarios by generating synthetic data that mimics real data without directly exposing sensitive information.
  - **Mechanisms**:
    - ✓ **Data Substitution**: GANs can replace real sensitive data (e.g., patient records) with synthetic equivalents in healthcare, while maintaining the statistical properties required for analysis.
    - ✓ **Differential Privacy**: By using differentially private GANs, it's possible to ensure that individual data points cannot be reverse-engineered from the synthetic data, preserving user privacy.

Apply Generative Adversarial Networks

Applications of GANs for Data Augmentation, Privacy, and Anonymity

- **Role in Privacy Preservation - Example**:

  ✓ **Healthcare**: GANs are used to generate synthetic patient data that resembles real patient data in terms of statistics, but without risking exposure of personally identifiable information (PII). This enables research without compromising patient privacy.

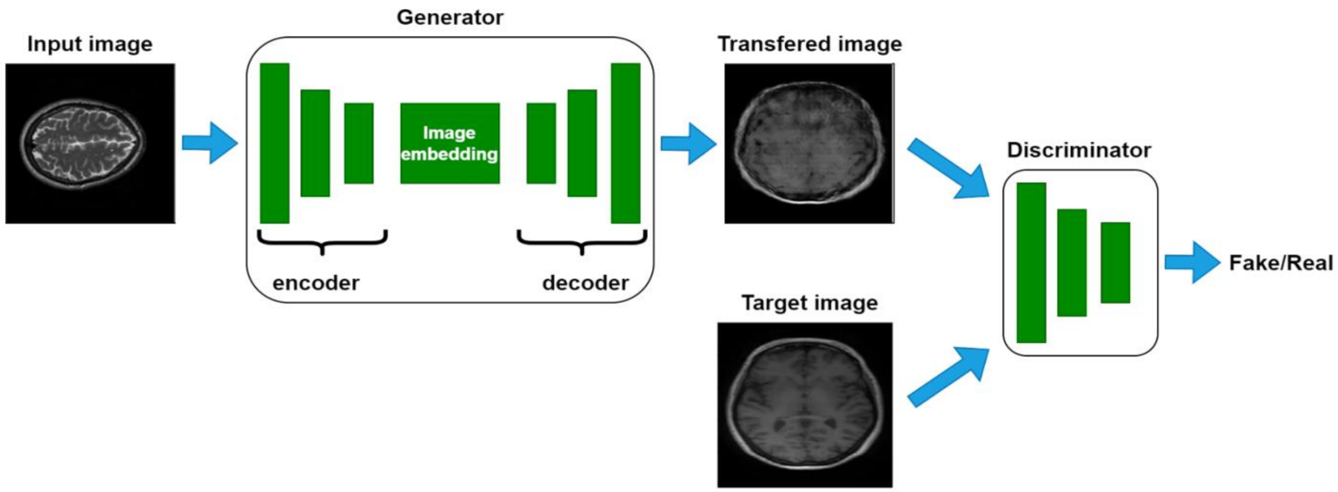

Image Source: https://www.mdpi.com/2079-9292/11/21/3470

Apply Generative Adversarial Networks

Applications of GANs for Data Augmentation, Privacy, and Anonymity

- **Anonymity in Sensitive Data Use Cases:**
  - ✓ GANs offer a unique advantage in anonymizing data while preserving its usability for analytics and model training, particularly for sensitive domains like healthcare or finance.
- **Application**:
  - ✓ In **financial data** settings, GANs can create synthetic transaction data that keeps patterns intact for model training while anonymizing the original data sources.

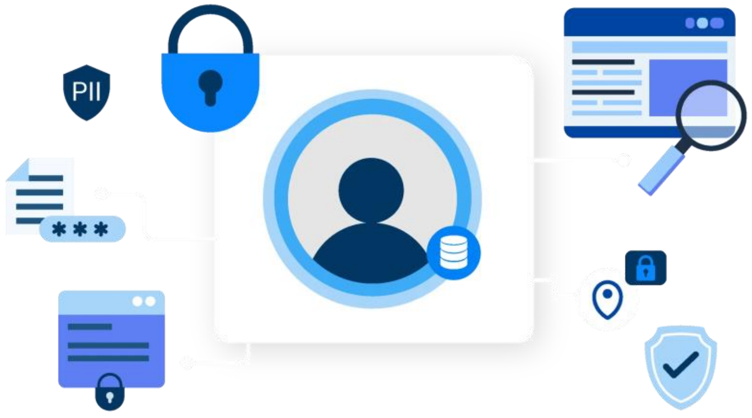

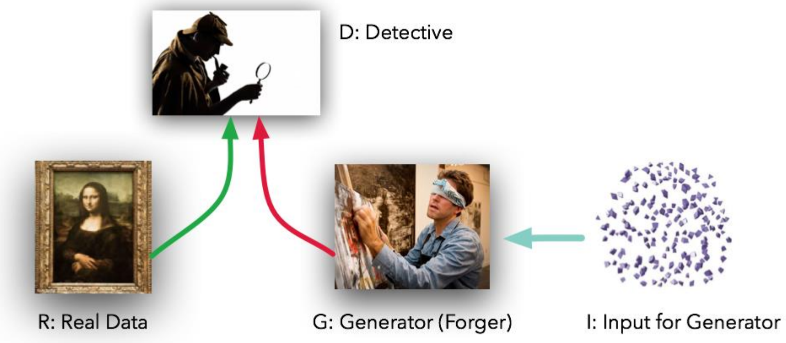

Image Source: https://becominghuman.ai/the-magic-of-generative-adversarial-network-gans-1c9d5da4a105

Apply Generative Adversarial Networks

Applications of GANs for Data Augmentation, Privacy, and Anonymity

- ✓ **Law enforcement**: GANs can generate anonymized video surveillance footage for training models without exposing individuals' identities.
- ✓ **Social Media**: GANs are used to anonymize user data, generating synthetic profiles or behavioral patterns for research without compromising individual privacy.

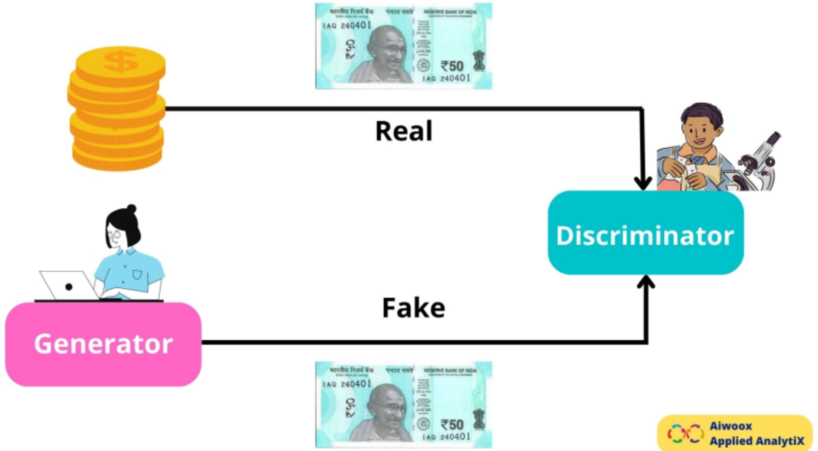

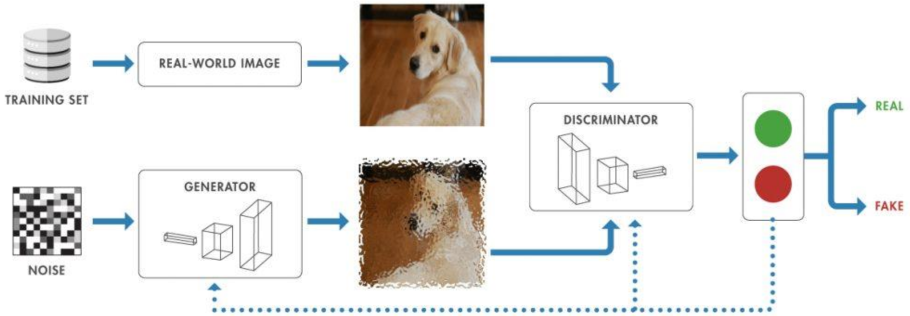

Image Source: https://blogs.mathworks.com/deep-learning/2021/12/02/synthetic-image-generation-using-gans/
Image Source: https://aiwoox.com/generative-adversarial-networks-gans/

Apply Generative Adversarial Networks

Applications of GANs for Data Augmentation, Privacy, and Anonymity

- **Case Studies (1)**
  - ✓ **Healthcare**:
    - **Data Augmentation**: GANs generate additional medical images (X-rays, MRIs, etc.) for training deep learning models in diagnosis.
    - **Privacy Preservation**: GANs create synthetic patient data to be used in research without violating HIPAA (Health Insurance Portability and Accountability Act) regulations.

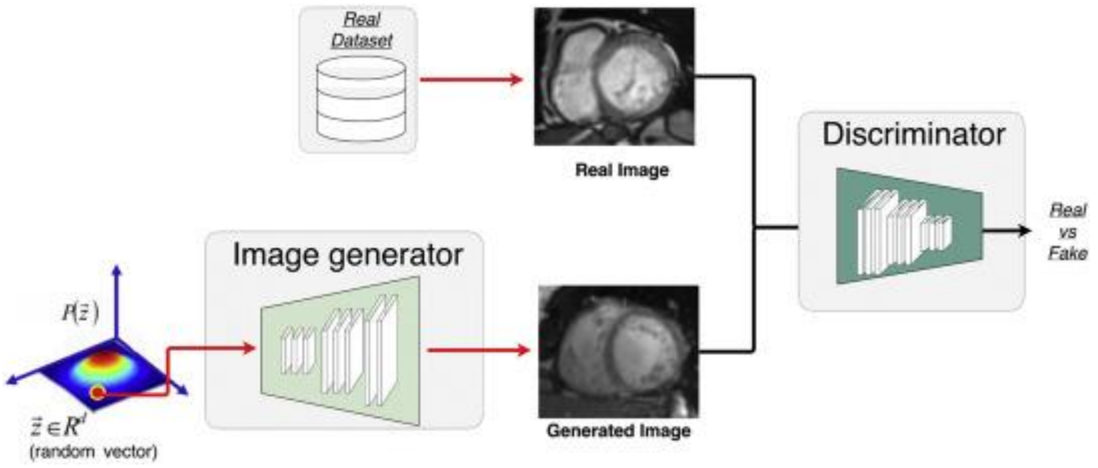

Image Source: https://www.sciencedirect.com/science/article/abs/pii/S0828282X21008606

Apply Generative Adversarial Networks

Applications of GANs for Data Augmentation, Privacy, and Anonymity

- **Case Studies (2)**
  - ✓ **Autonomous Driving**:
    - **Data Augmentation**: GANs simulate diverse driving scenarios, weather conditions, and road hazards to augment real-world driving data, helping autonomous vehicles learn to handle varied situations.
    - **Privacy**: In smart city applications, GANs can anonymize pedestrian and vehicle data used in traffic monitoring systems to preserve privacy while still supporting data-driven decisions.

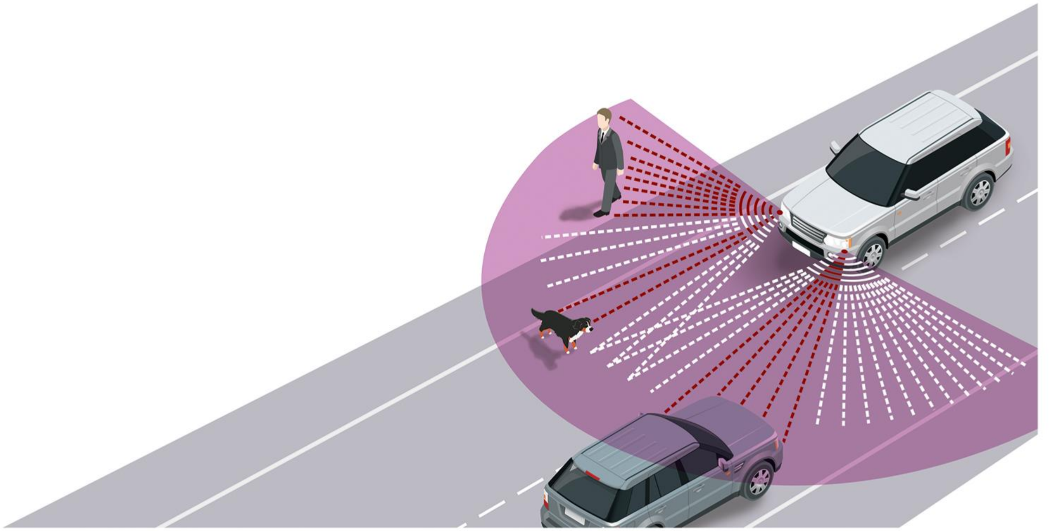

Image Source: https://epc-co.com/epc/jp/epc%E3%81%AB%E3%81%A4%E3%81%84%E3%81%A6/gan%E3%81%AE%E8%A9%B1/post/13749/gan-technology-for-the-connected-car

Apply Generative Adversarial Networks

Applications of GANs for Data Augmentation, Privacy, and Anonymity

- **Improving Downstream AI Models with GAN-Generated Data**

  ✓ GANs excel at generating synthetic data that enhances the diversity of training datasets, particularly in scenarios where real-world data is limited, imbalanced, or expensive to collect.

Apply Generative Adversarial Networks

Improving Downstream AI Models with GAN-Generated Data

- **Key Benefits:**
  - ✓ **Increased Variability**: GANs can generate new examples by capturing the underlying distribution of real data and introducing variations such as new angles, lighting conditions, or slight transformations, leading to a more robust training set.
  - ✓ **Handling Imbalanced Datasets**: GANs are particularly useful for balancing datasets. For example, in medical diagnosis, rare disease images can be generated to match the number of healthy samples, thus mitigating class imbalance.
  - ✓ **Edge Cases**: GANs can simulate rare or dangerous scenarios that may be difficult to capture in the real world. For instance, in autonomous driving, generating rare events (like accidents or unusual lighting conditions) can train AI models to handle such cases better.

Apply Generative Adversarial Networks

Improving Downstream AI Models with GAN-Generated Data

- **Key Benefits-Ex:**

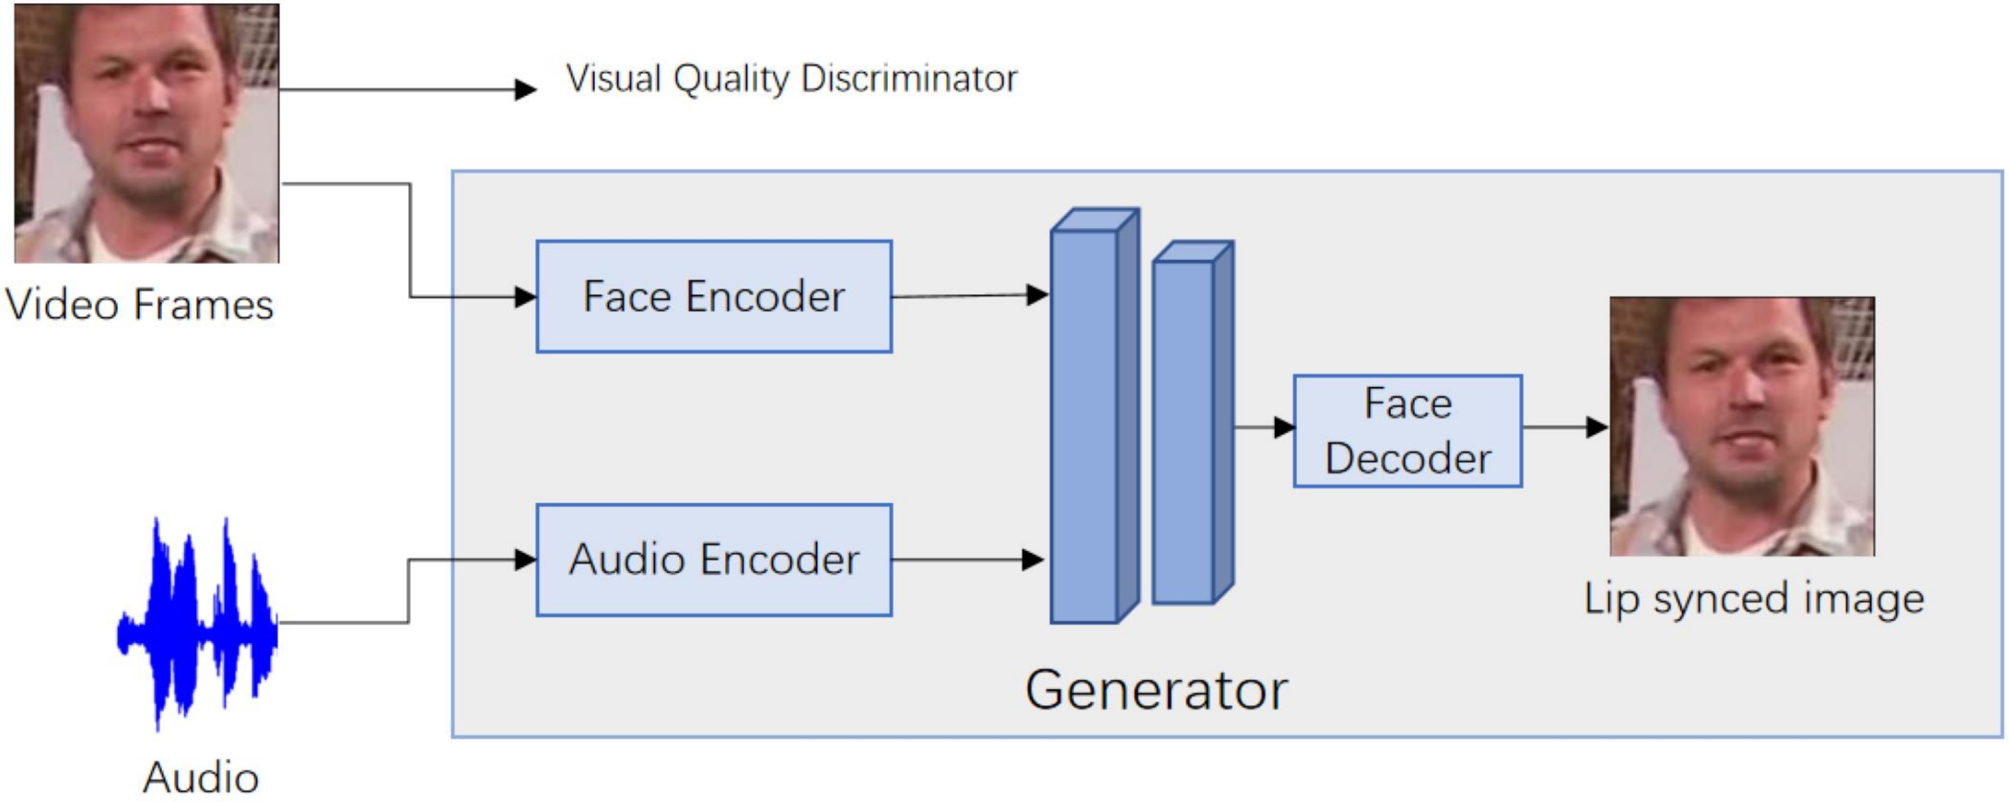

Apply Generative Adversarial Networks

Improving Downstream AI Models with GAN-Generated Data

- **Enhancing Performance of AI Models with GAN Data**

  ✓ By improving the quality and diversity of data, GAN-generated datasets can significantly boost the performance of downstream AI models.

  - **Key Improvements**:

    1. **Generalization**: Models trained with both real and GAN-generated data are better equipped to generalize to new, unseen data. This is because GANs introduce subtle variations that help the model learn a wider range of patterns.

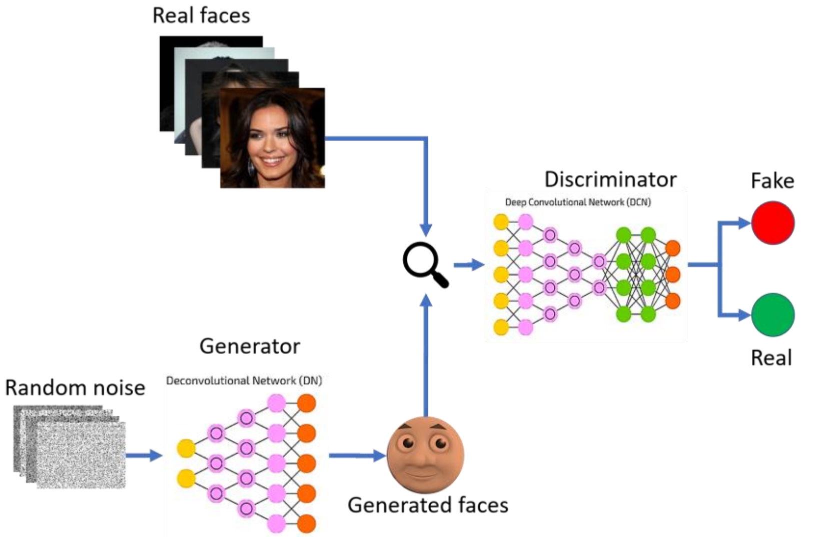

Image Source: https://www.spindox.it/en/generative-adversarial-neural-networks/

Apply Generative Adversarial Networks

Improving Downstream AI Models with GAN-Generated Data

- **Enhancing Performance of AI Models with GAN Data**
  - ✓ **Key Improvements**:
    - **Overfitting Prevention**: With limited datasets, models can overfit, learning the noise of the training data rather than its general patterns.
    - GANs provide additional synthetic data that reduces the chances of overfitting, leading to better performance on test data.

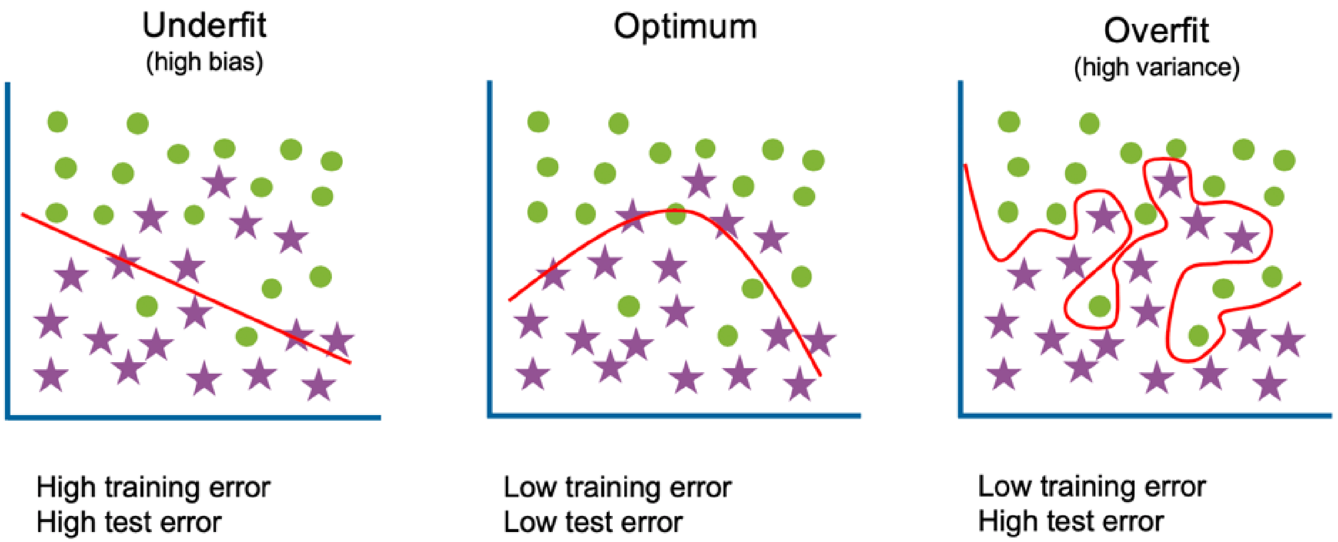

Image Source: https://www.linkedin.com/pulse/overfitting-underfitting-machine-learning-ml-concepts-com

Apply Generative Adversarial Networks

Improving Downstream AI Models with GAN-Generated Data

- **Enhancing Performance of AI Models with GAN Data**

✓ **Key Improvements**:

- **Model Accuracy**: GANs help improve the precision and recall of AI models by ensuring they are trained on a diverse set of data points.

- This is especially useful in fields like object detection and classification, where diverse inputs improve recognition across different conditions.

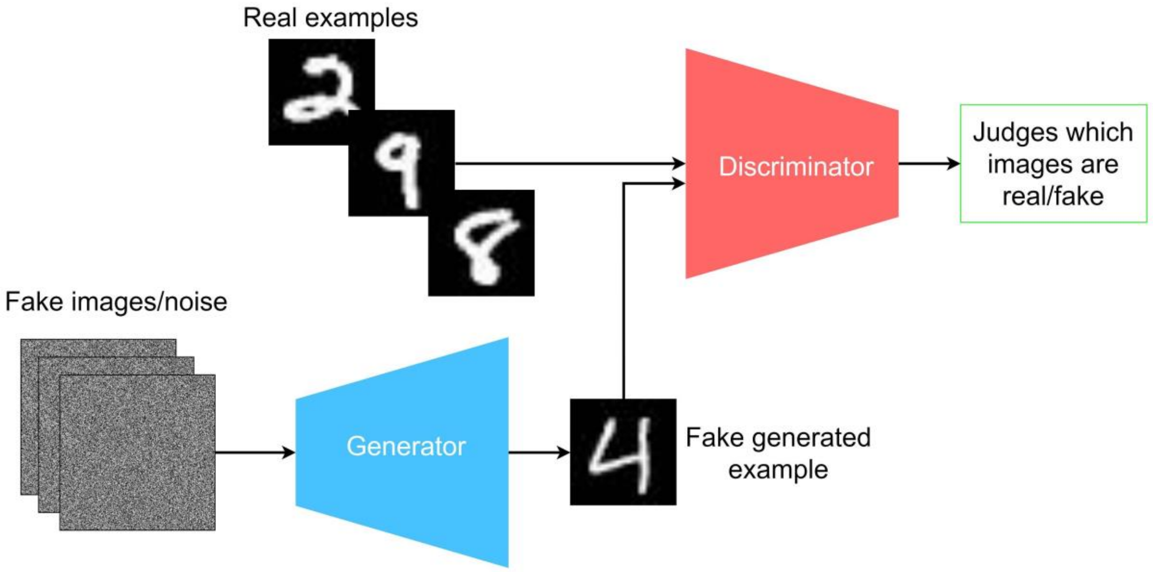

Image Source: https://developer.ibm.com/articles/generative-adversarial-networks-explained/

Apply Generative Adversarial Networks

Improving Downstream AI Models with GAN-Generated Data

- **Case Study Example: GAN-Generated Data in Object Detection**

  ✓ **Scenario**: In object detection, models often need large amounts of varied data to perform well across different environments and object appearances.

Use Case:

✓ **Problem**: A company developing AI models for autonomous drones faces difficulty in obtaining large amounts of labeled training data in varied weather, lighting, and terrain conditions. Traditional data collection methods are costly and time-consuming.

✓ **Solution**: GANs are used to augment the real-world data by generating synthetic images of the same objects (e.g., vehicles, trees, people) under different conditions—rain, snow, fog, and at various times of day.

Apply Generative Adversarial Networks

Improving Downstream AI Models with GAN-Generated Data

- **Case Study Example: GAN-Generated Data in Object Detection**
  - **Outcome**:
    - ✓ The object detection model's accuracy increases significantly due to exposure to a wider variety of training images.
    - ✓ The model becomes more robust to real-world situations where lighting or weather conditions are different from those in the original dataset.
    - ✓ The GAN-augmented training data reduces the need for expensive and time-consuming manual data collection.

Apply Generative Adversarial Networks

Improving Downstream AI Models with GAN-Generated Data

- **Case Study Example: GAN-Generated Data in Object Detection**
  - **Results**:

✓ **Performance Improvement**: The object detection model shows a 20% improvement in recall and 15% improvement in precision when tested on real-world footage after training with GAN-augmented data.

✓ **Cost Efficiency**: The use of GAN-generated data reduced the need for additional data collection by 30%, saving both time and resources.

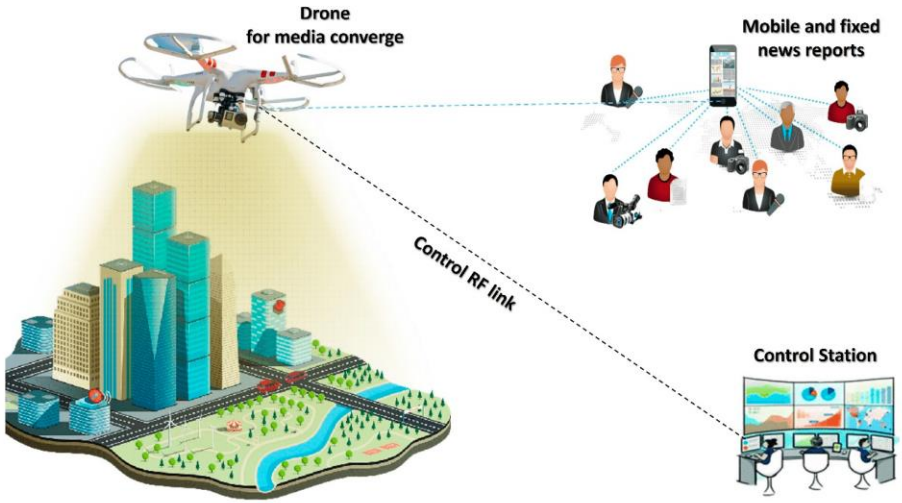

Image Source: https://www.mdpi.com/2071-1050/14/7/3758

Apply Generative Adversarial Networks

Image-to-Image Translation

- ✓ Image-to-image translation is a subfield of computer vision where the goal is to map an input image from one domain to a corresponding output image in another domain.
- ✓ It involves learning the transformation between the two domains through deep learning models, particularly generative models like Generative Adversarial Networks (GANs).

Apply Generative Adversarial Networks

Image-to-Image Translation

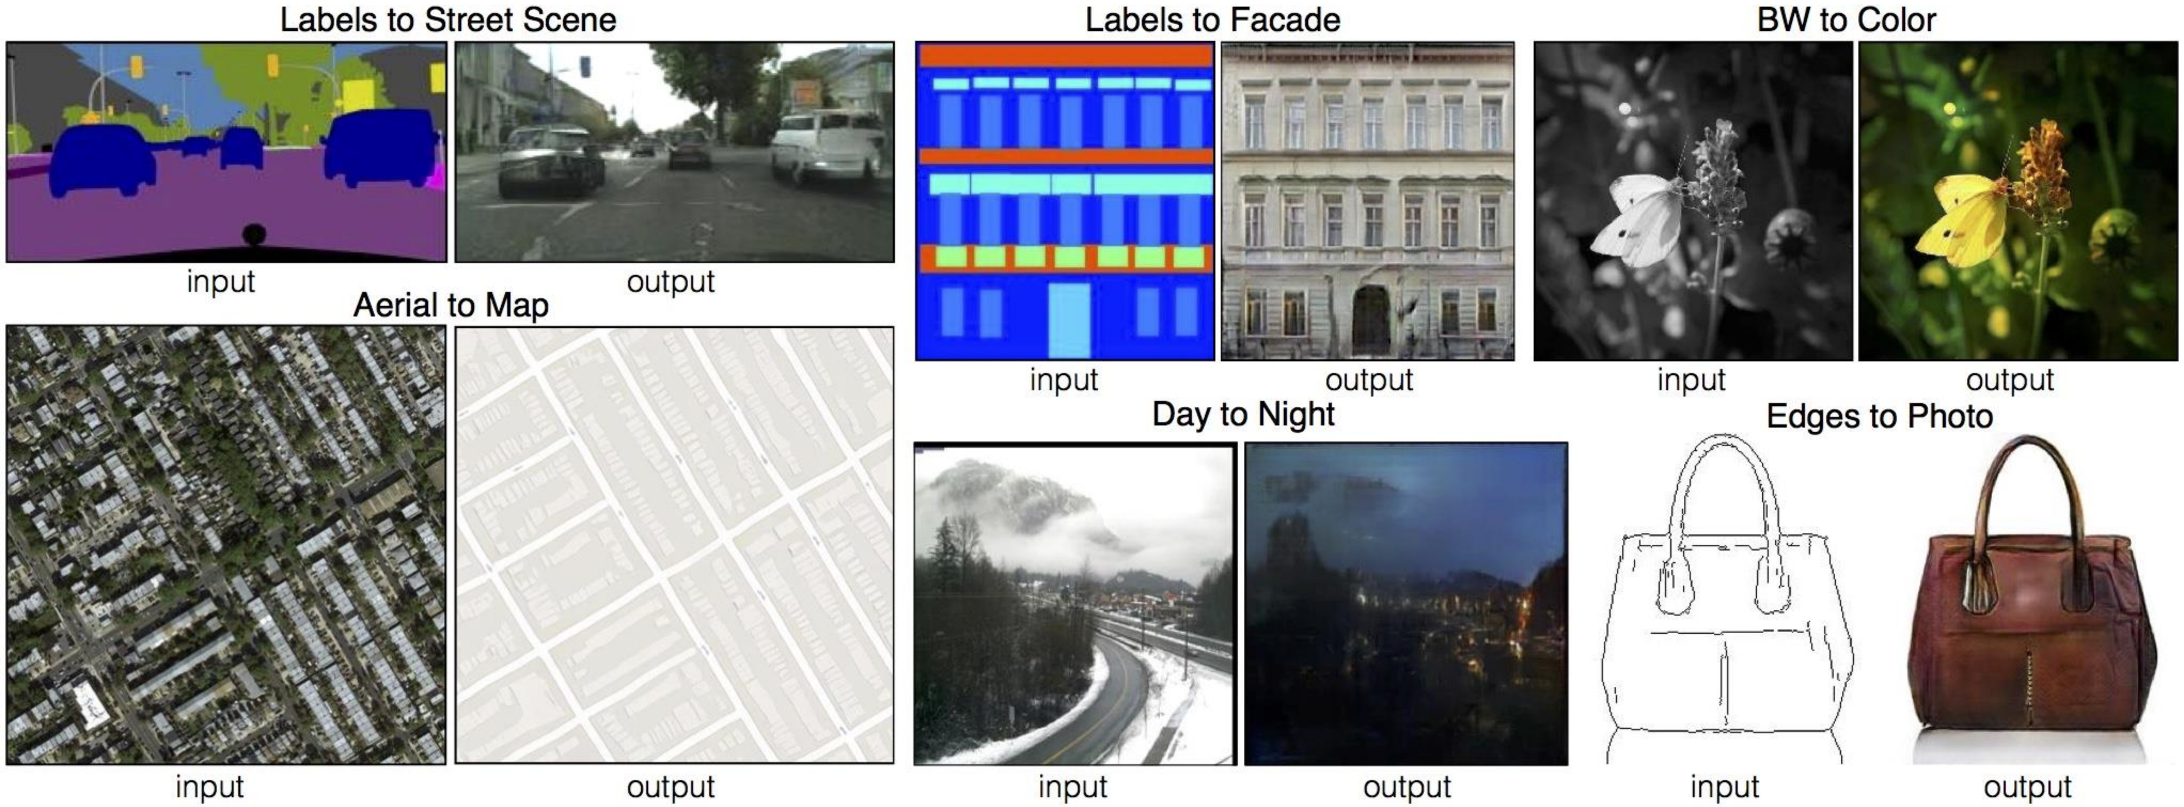

Image Source: https://phillipi.github.io/pix2pix/

Apply Generative Adversarial Networks

Image-to-Image Translation

✓ **Key Characteristics:**

- The task is to modify certain aspects of the input image (style, lighting, or content) while maintaining some consistent structure across domains.
- It typically relies on paired or unpaired datasets depending on the specific model used.
- Requires both the generator (to create new images) and the discriminator (to distinguish real from fake images) for effective learning.

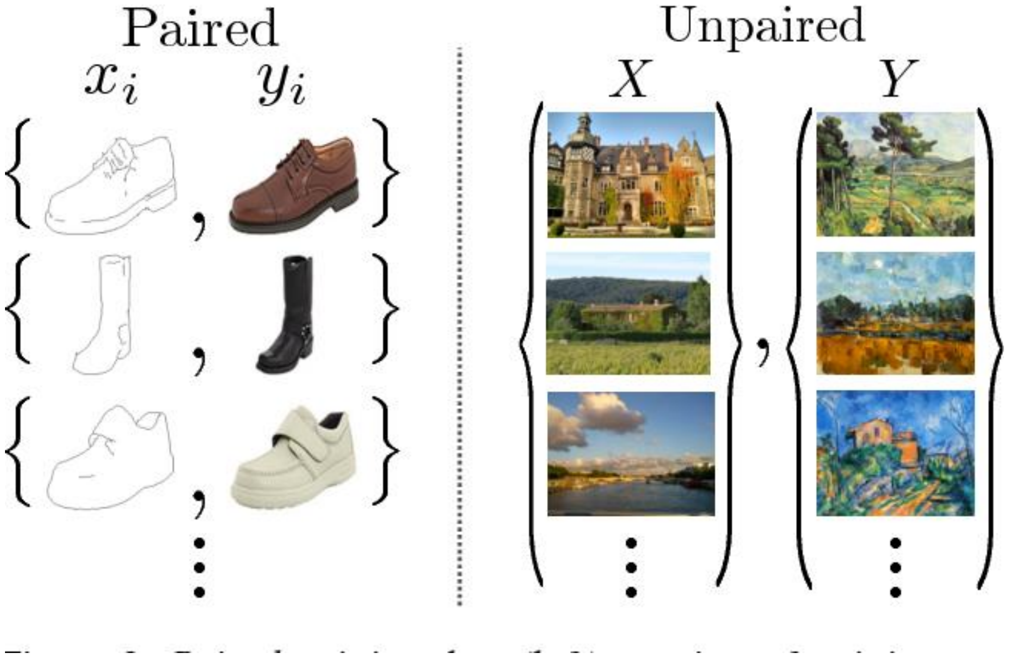

Image Source: https://medium.com/abacus-ai/gans-for-data-augmentation-21a69de6c60b

Apply Generative Adversarial Networks

Image-to-Image Translation

- **Popular Applications:**
  - **Style Transfer:**
    - ✓ Style transfer involves applying the visual appearance (style) of one image onto the content of another.
    - ✓ **Example:** Transforming a landscape photo into the style of a famous painting (e.g., turning a natural photograph into the style of Van Gogh's "Starry Night").
    - ✓ Used in artistic applications, content creation, and photography enhancements.

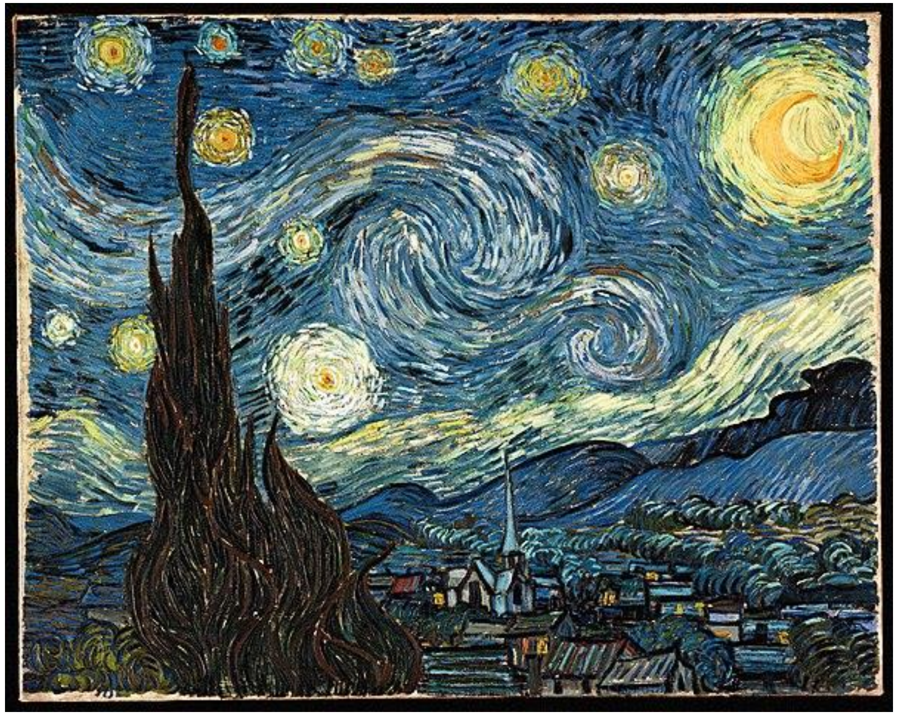

Image Source: https://commons.wikimedia.org/wiki/File:VanGogh-starry_night_edit

Apply Generative Adversarial Networks

Image-to-Image Translation

- **Popular Applications:**
  - **Domain Adaptation:**
    - ✓ This refers to transforming images from one domain to another to reduce the domain gap, often to improve model performance in scenarios where labeled data is scarce.
    - ✓ **Example:** Converting synthetic data to real-world data for training AI models in applications like autonomous driving or facial recognition.
    - ✓ Helps overcome the challenges of training on data that differs significantly from the target domain.

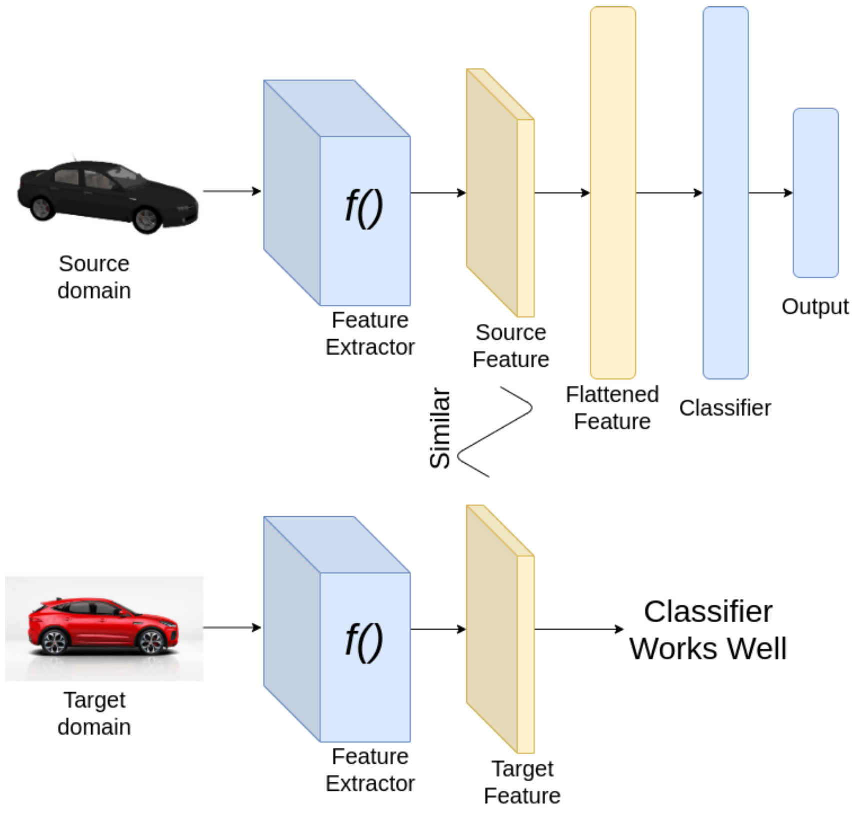

Apply Generative Adversarial Networks

Image-to-Image Translation

- **Popular Applications:**
  - **Medical Imaging:**
    - ✓ Image-to-image translation in medical imaging often helps in converting one imaging modality to another (e.g., translating MRI scans to CT scans or vice versa).
    - ✓ **Example:** Enhancing the quality of low-resolution medical images or converting a medical image to a different format to aid diagnosis.
    - ✓ This technique aids in diagnosis, treatment planning, and even reducing radiation exposure by generating synthetic scans.

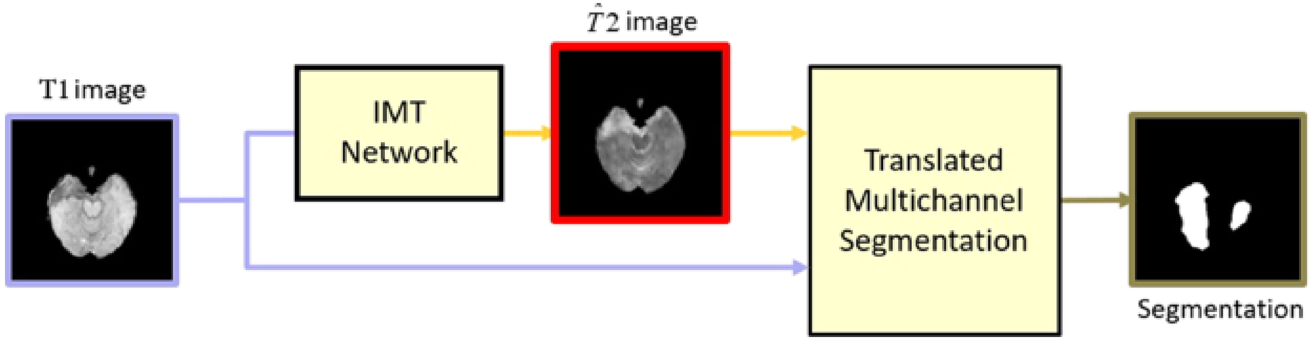

Apply Generative Adversarial Networks

Paired Image-to-Image Translation

- **Paired image-to-image translation**, each image in the source domain has a corresponding image in the target domain, making it easier to learn a mapping between the two.
- This approach is particularly useful when datasets are available with clear, one-to-one relationships between input and output images, such as converting a grayscale image to a color image or transforming a blueprint into a 3D model.

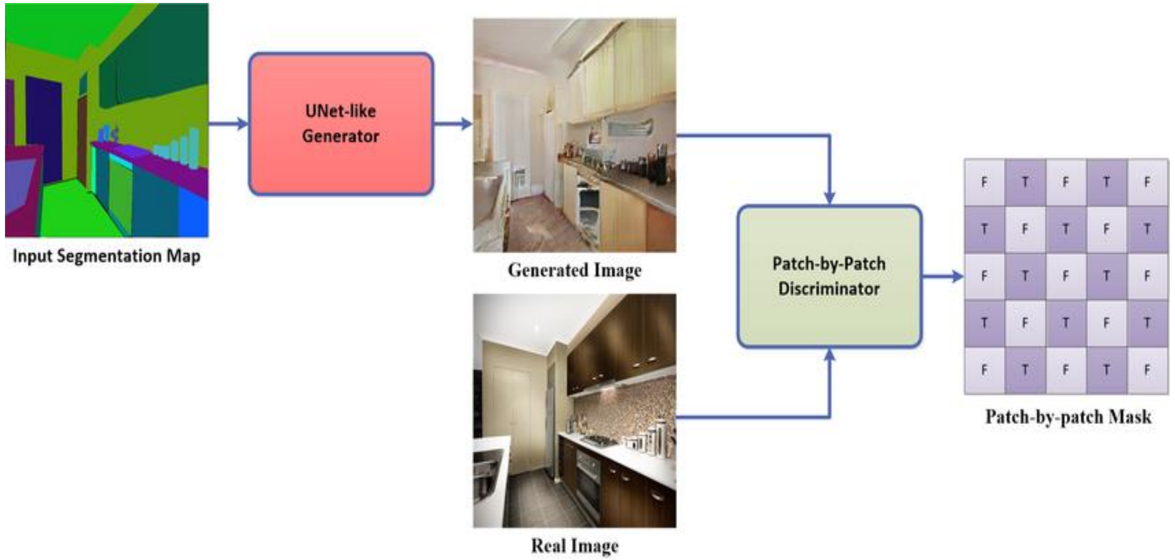

Image Source: https://www.researchgate.net/figure/Overview-of-the-proposed-paired-Image-to-Image-translation-technique_fig1_364973747

Apply Generative Adversarial Networks

Paired Image-to-Image Translation

- Pix2Pix, one of the most popular models for paired image-to-image translation, leverages this pairing to learn how to generate the target image given a corresponding input image.

- It uses a GAN architecture that includes a **Generator** and a **Discriminator**, but it enhances the traditional GAN model by specifically training on paired data, ensuring higher quality and consistency in the output.

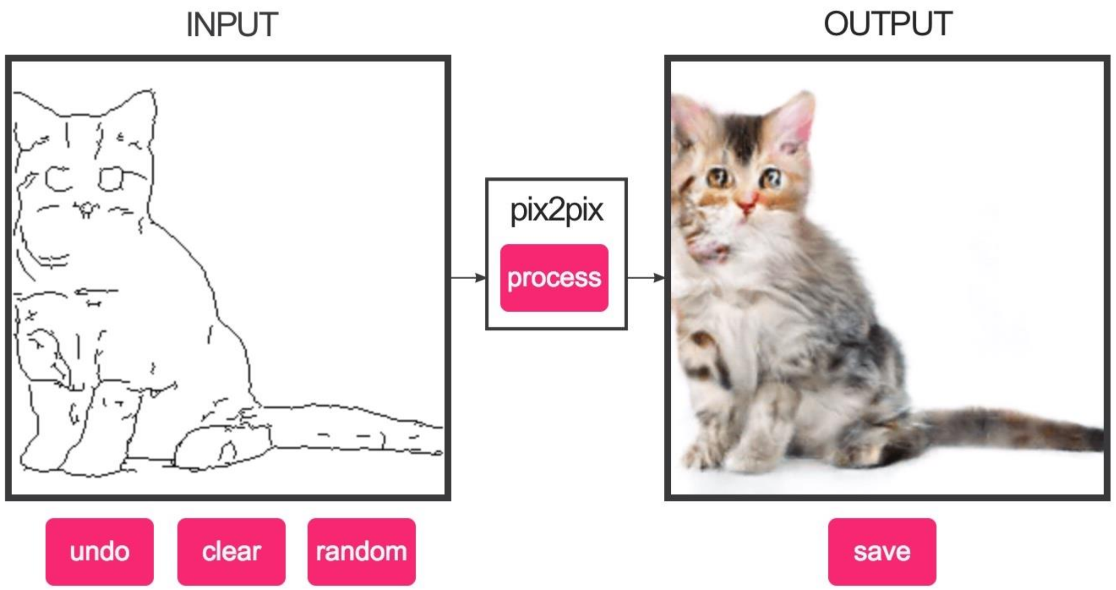

Image Source: https://phillipi.github.io/pix2pix/

Apply Generative Adversarial Networks

Paired Image-to-Image Translation

- **How Pix2Pix Works: Paired Data Example (Satellite Images ↔ Map Routes)**

  Pix2Pix is particularly well-suited for problems where paired data is available, such as transforming **satellite images into map routes**:

  1. **Data Setup**:

     In this example, each satellite image (source) is paired with its corresponding map (target). These paired images are fed into the model, and the goal is to generate a realistic map from a satellite image.

Image Source: https://www.mdpi.com/1424-8220/20/11/3119

Apply Generative Adversarial Networks

Paired Image-to-Image Translation

- **How Pix2Pix Works: Paired Data Example (Satellite Images ↔ Map Routes)**

2. **Generator**:
   - ✓ The generator in Pix2Pix is designed to take a satellite image as input and output a corresponding map image.
   - ✓ It uses an advanced **U-Net** architecture, which allows the generator to produce high-quality images by capturing both low-level and high-level features.

3. **Discriminator**:
   - ✓ The discriminator works by trying to distinguish between real map images and the generated ones.
   - ✓ Pix2Pix uses a **PatchGAN discriminator**, which evaluates image quality on a patch-by-patch basis instead of evaluating the entire image.
   - ✓ This helps focus on local details, making the generated images more realistic.

Apply Generative Adversarial Networks

Paired Image-to-Image Translation

- **How Pix2Pix Works: Paired Data Example (Satellite Images ↔ Map Routes)**

5. **Training Process**:
   - ✓ The generator and discriminator compete in a min-max game. The generator tries to fool the discriminator by producing high-quality maps that are indistinguishable from real ones, while the discriminator tries to correctly classify real versus generated maps.
   - ✓ Over time, the generator improves, and the model learns to generate highly accurate map routes from satellite images.

6. **Loss Function**:
   - ✓ Pix2Pix uses a combination of **adversarial loss** (to fool the discriminator) and **L1 loss** (to encourage the generated image to be close to the ground-truth paired image).
   - ✓ The L1 loss ensures that the generator produces outputs that not only look realistic but also match the paired target image in terms of pixel values.

Apply Generative Adversarial Networks

Paired Image-to-Image Translation

- **Architectures: Advanced U-Net Generator and PatchGAN Discriminator**
  - **U-Net Generator**:
    - ✓ U-Net is a fully convolutional neural network with a **"U" shape**. It has **encoder-decoder architecture**:

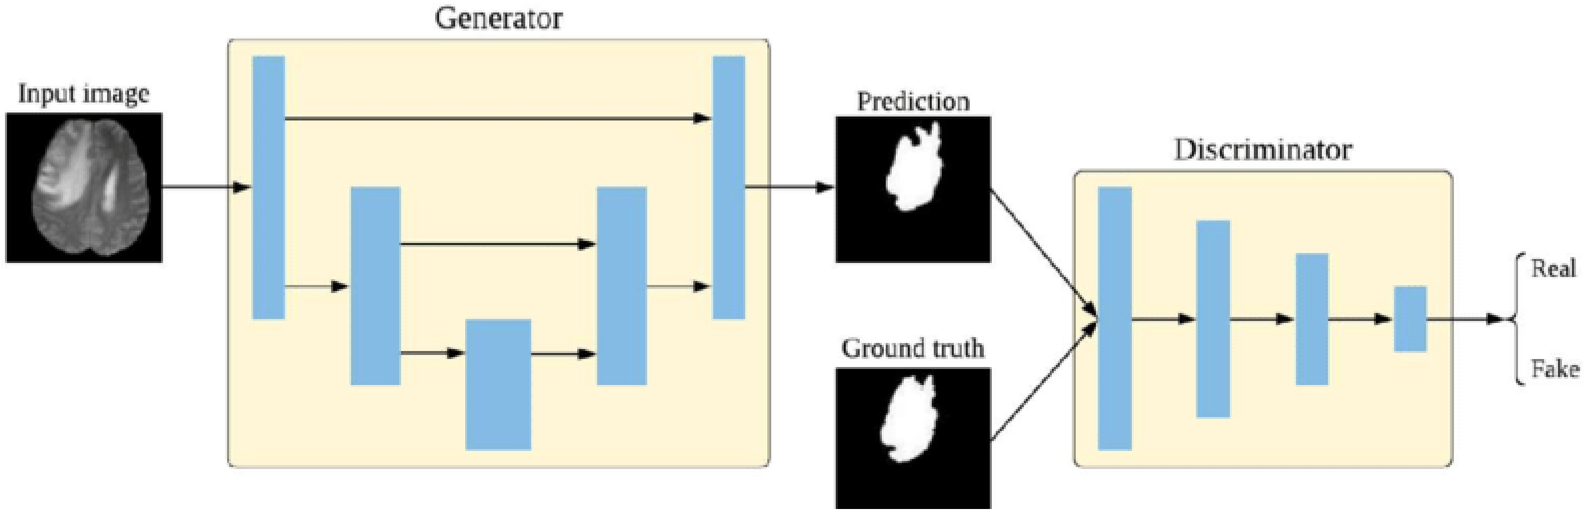

Image Source: https://www.researchgate.net/figure/Simplified-schematic-of-U-net-based-GAN-The-generator-synthesizes-predictions-for-the_fig5_352111620

Apply Generative Adversarial Networks

Paired Image-to-Image Translation

- **Architectures: Advanced U-Net Generator and PatchGAN Discriminator**

2. PatchGAN Discriminator:

✓ Instead of evaluating the entire image at once, PatchGAN works by dividing the image into small **patches** (e.g., 70x70 pixels) and evaluating whether each patch is real or fake.

✓ This ensures that the discriminator focuses on local details and textures rather than the overall image.

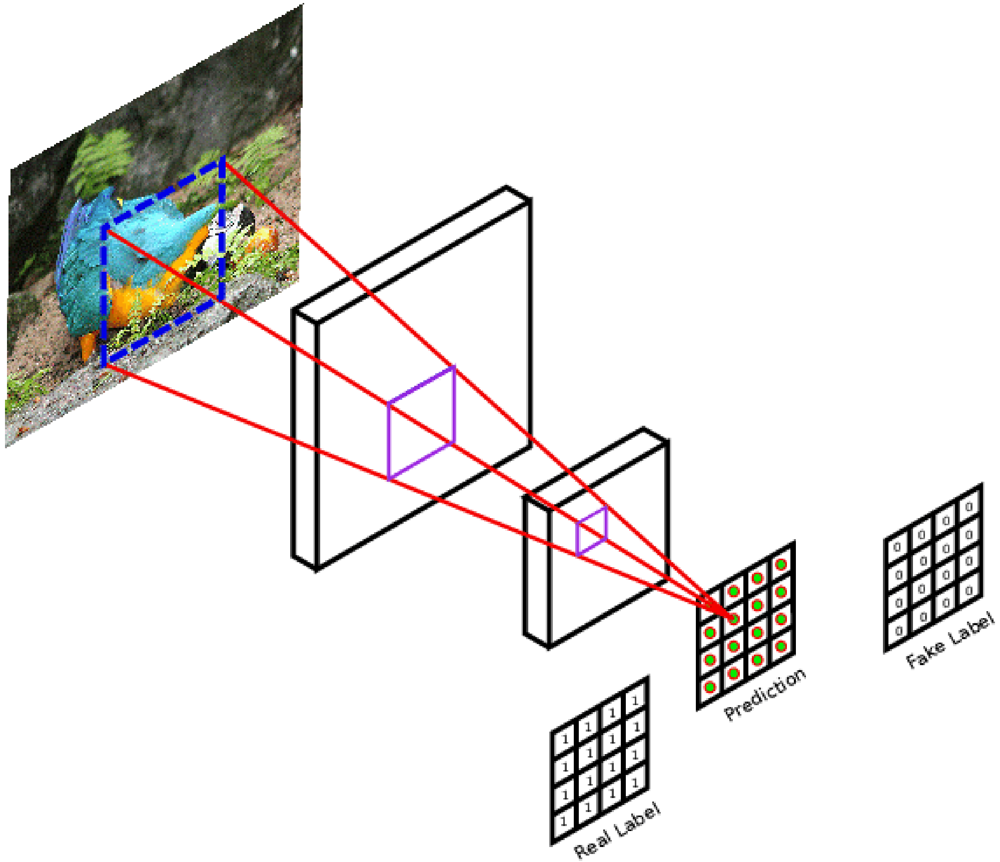

https://www.researchgate.net/figure/PatchGAN-discriminator-Each-value-of-the-output-matrix-represents-the-probability-of_fig1_323904616

Apply Generative Adversarial Networks

Paired Image-to-Image Translation

- **Satellite Images ↔ Map Routes**
  - ✓ **Urban planning** or **Geographical analysis**, the goal is to automatically generate detailed and accurate map routes from satellite images.
  - ✓ Using Pix2Pix, a model can be trained on paired satellite and map images to learn this mapping and produce maps without manual intervention.

Image Source: https://www.geospatialworld.net/blogs/gis-in-urban-planning/

Apply Generative Adversarial Networks

Implementation of Pix2Pix

Step 1: Dataset Preparation

- ✓ Collect paired image data (e.g., satellite images paired with corresponding map routes).
- ✓ Split the dataset into training, validation, and test sets.
- ✓ Preprocess the images (e.g., resizing, normalization) for input into the neural network.

Step 2: Define the U-Net Generator

- ✓ Construct the **U-Net architecture**, starting with the encoder and decoder.
- ✓ Use **convolutional layers** in the encoder to extract features and **transpose convolutions** in the decoder to reconstruct the output.
- ✓ Implement **skip connections** between corresponding encoder and decoder layers to preserve spatial information and details.

Apply Generative Adversarial Networks

Implementation of Pix2Pix

Step 3: Define the PatchGAN Discriminator

- ✓ Build the **PatchGAN discriminator** to operate on small image patches (e.g., 70x70).
- ✓ This discriminator classifies each patch as real or fake, improving local feature generation.

Step 4: Define the Loss Functions

- ✓ Implement the **adversarial loss** to ensure the generator produces realistic images that fool the discriminator.
- ✓ Add the **L1 loss** to minimize the pixel difference between the generated image and the actual target image (ground truth), enhancing image fidelity.

Apply Generative Adversarial Networks

Implementation of Pix2Pix

Step 5: Train the Model

- ✓ Train the model in a loop where:
  - ✓ The generator creates a target domain image (e.g., map from satellite image).
  - ✓ The discriminator evaluates both the real and generated images.
  - ✓ Losses are backpropagated to improve the performance of both the generator and discriminator.
- ✓ Use the validation set to monitor the model's performance and prevent overfitting.

Apply Generative Adversarial Networks

Implementation of Pix2Pix

Step 7: Fine-Tuning and Optimization

- ✓ Adjust hyperparameters (e.g., learning rate, batch size) to optimize performance.
- ✓ Perform additional training cycles or **fine-tune** the model based on specific applications.

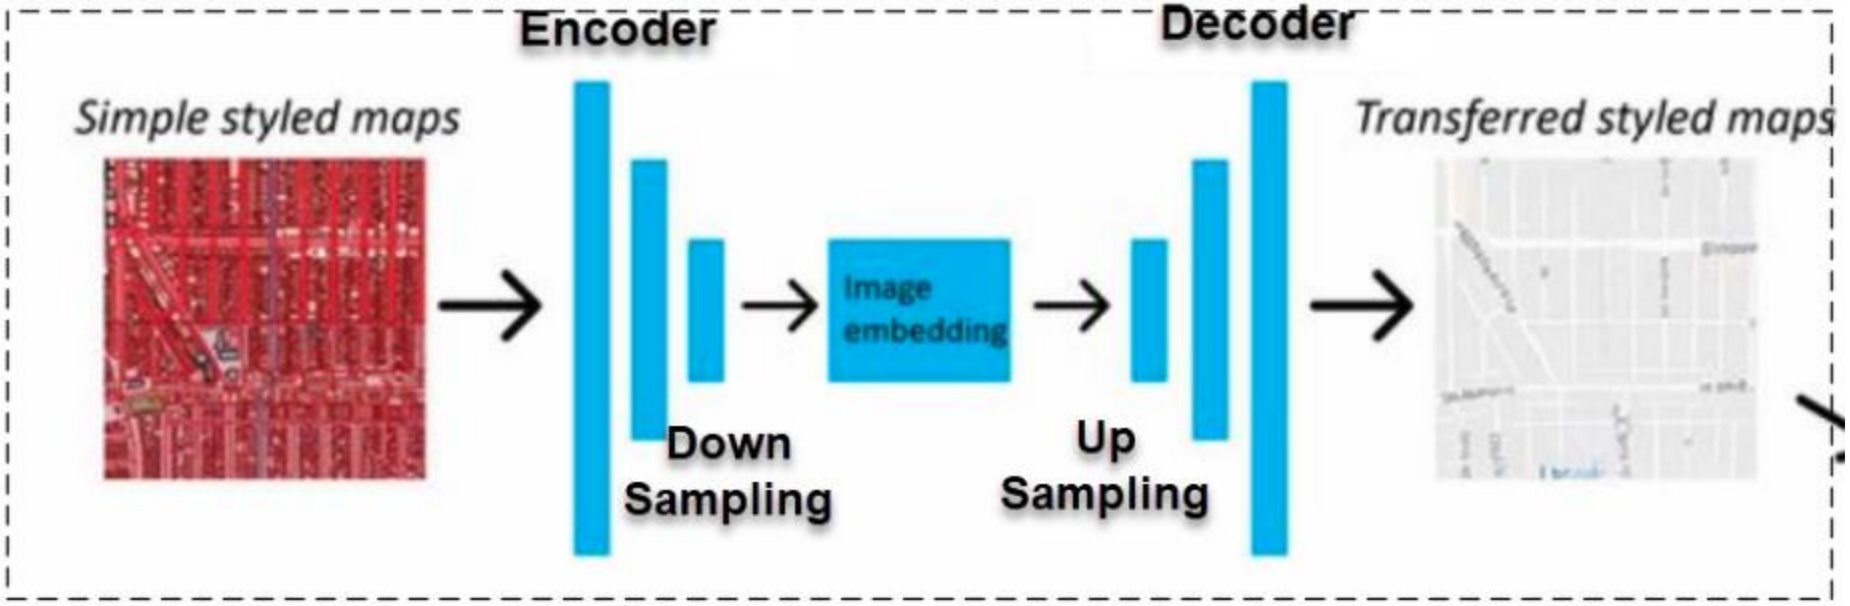

Image Source: https://phamdinhkhanh.github.io/2020/11/13/pix2pixGAN.html

Apply Generative Adversarial Networks

Adapting Satellite Images into Maps (and vice versa)

- **Use Case: Satellite to Map Translation**
  - ✓ The trained Pix2Pix model can convert satellite images into detailed map routes for applications like urban planning, disaster management, or autonomous navigation.
- **Reverse Translation: Map to Satellite**
  - ✓ The model can also be trained in the reverse direction, generating satellite-like images from map data, aiding in visualization and simulation tasks.
- **Real-World Impact**
  - ✓ Helps in creating efficient tools for geospatial data analysis, traffic systems, and infrastructure development, where automated image-to-image translation can save time and resources.

Apply Generative Adversarial Networks

Architecture Walkthrough: U-Net and PatchGAN

- **U-Net Generator**
  - ✓ **Encoder**: The encoder part of the U-Net progressively downsamples the input image (e.g., a satellite image) to capture abstract features and reduce dimensionality.
  - ✓ **Decoder**: The decoder upsamples the image back to its original resolution while using **skip connections** to maintain spatial consistency and detailed features.
  - ✓ **Skip Connections**: These connections allow the model to bypass certain layers, transferring lower-level features directly to the decoder, improving the quality of the generated images.

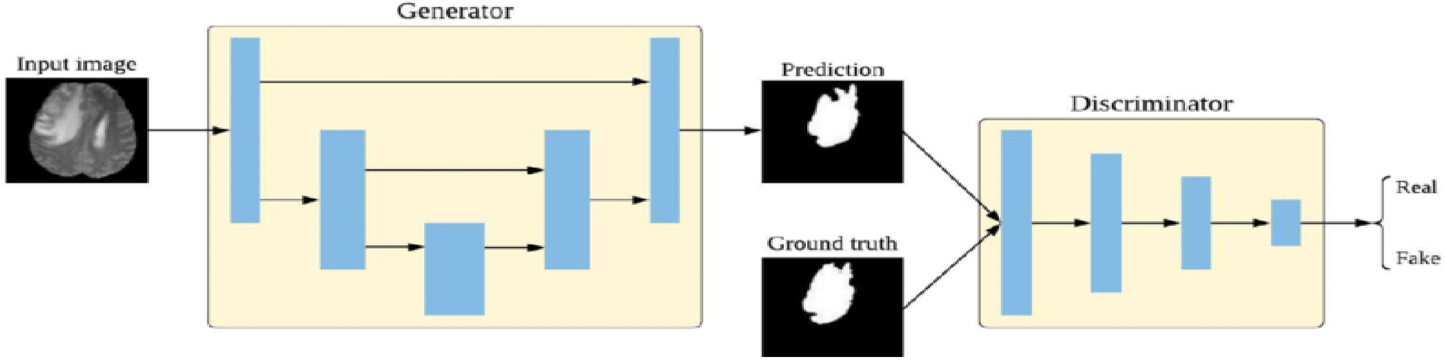

Image Source: https://www.researchgate.net/figure/Simplified-schematic-of-U-net-based-GAN-The-generator-synthesizes-predictions-for-the_fig5_352111620

Apply Generative Adversarial Networks

Architecture Walkthrough: U-Net and PatchGAN

- **PatchGAN Discriminator**
  - ✓ **Patch-Based Discrimination**: Instead of evaluating the entire image, the PatchGAN discriminator evaluates individual patches (70x70 pixels) to ensure that local features and textures are realistic.
  - ✓ **Global Consistency via Local Evaluations**: By focusing on smaller regions, PatchGAN improves the generator's ability to produce fine details like edges, textures, and small objects.

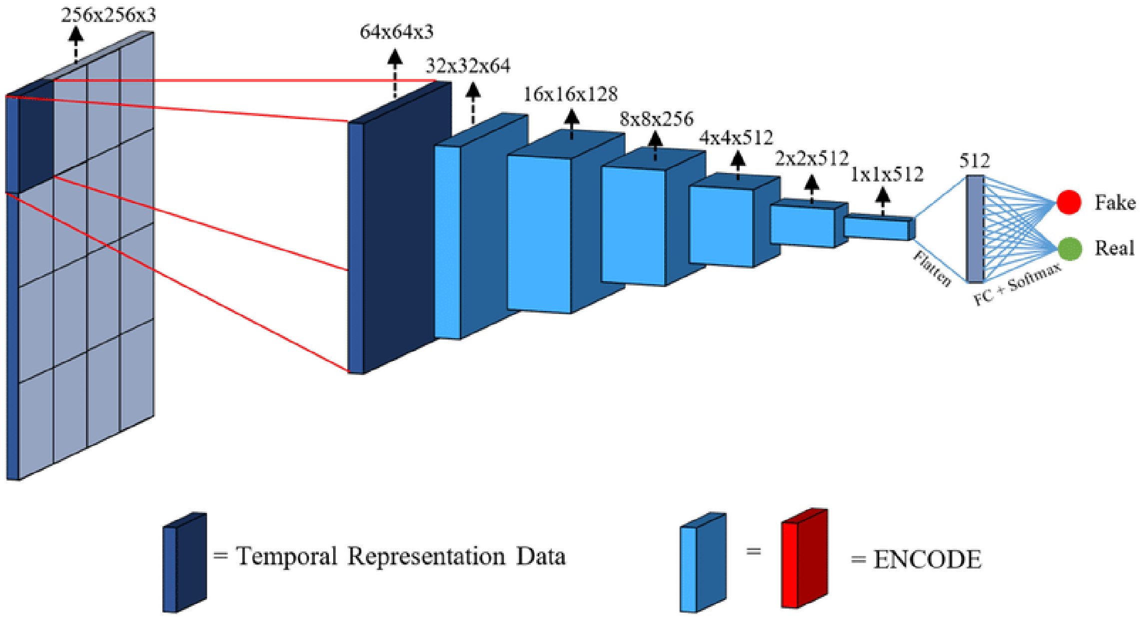

Image Source: https://www.researchgate.net/figure/The-PatchGAN-structure-in-the-discriminator-architecture_fig5_339832261

Image-to-Image Unpaired Translation

Apply Generative Adversarial Networks

Image-to-Image Unpaired Translation

- ✓ Unpaired Image-to-Image Translation is a technique in which the source and target images do not have one-to-one correspondence.
- ✓ In simpler terms, there is no requirement for direct pairing between input and output images in the training data, unlike paired image-to-image translation.
- ✓ The model learns to map a domain (e.g., horses) to another domain (e.g., zebras) without being provided with exact before-and-after images of horses turning into zebras.

Apply Generative Adversarial Networks

Image-to-Image Unpaired Translation

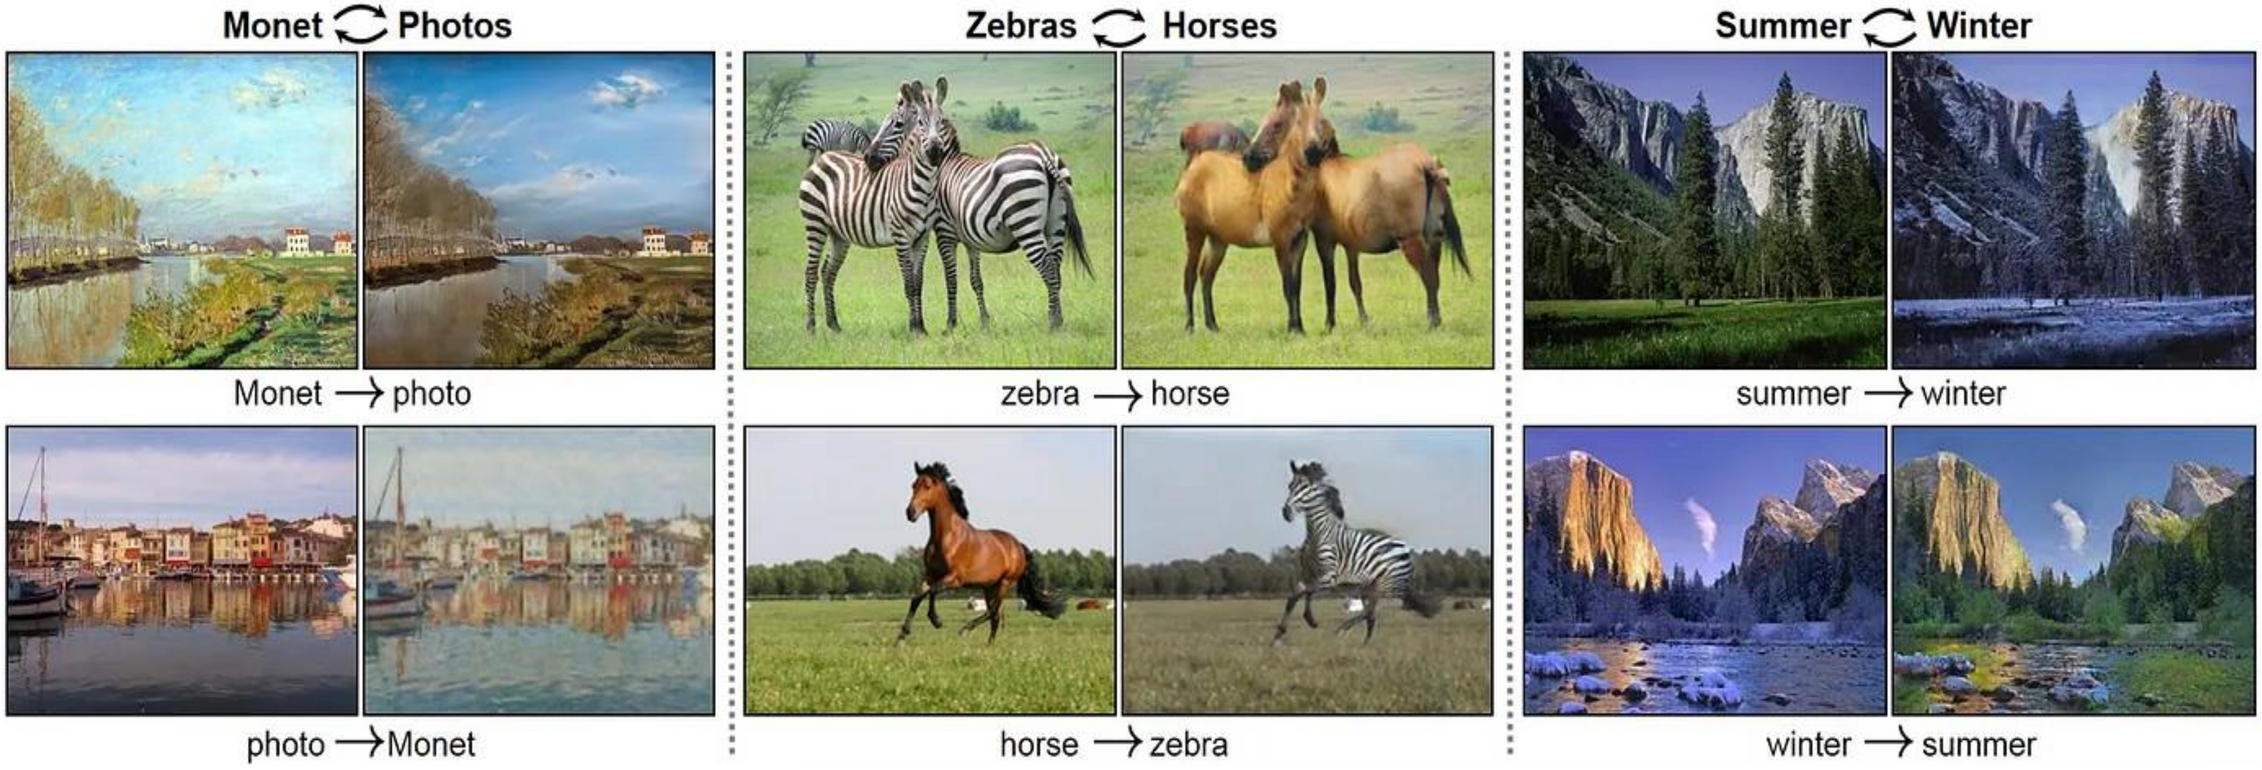

Image Source: https://towardsdatascience.com/cyclegan-how-machine-learning-learns-unpaired-image-to-image-translation-3fa8d9a6aa1d

Apply Generative Adversarial Networks

key differences between Paired and Unpaired Translation

Aspect	Paired Translation	Unpaired Translation
Input Data	Requires aligned image pairs (e.g., an image and its corresponding transformed version like a day-to-night pair)	No aligned image pairs are provided.
Learning Approach	The model learns the transformation directly from the paired data.	The model learns domain characteristics independently and maps between them.
Mapping Difficulty	Easier to map the relationship between domains since pairs provide clear guidance.	More challenging as the model must retain essential characteristics without direct correspondence.
Application	Suited for cases where pairs of images in different domains are readily available.	Suitable for cases where aligned pairs are not available, relying on domain-specific features for mapping.

Apply Generative Adversarial Networks

key differences between Paired and Unpaired Translation

Example:

- ✓ **Paired Translation:** Translating an aerial image to a corresponding map, where both the original and transformed versions are aligned.

- ✓ **Unpaired Translation:** Translating photos of horses to zebras, where the model doesn't have matching images to directly learn from.

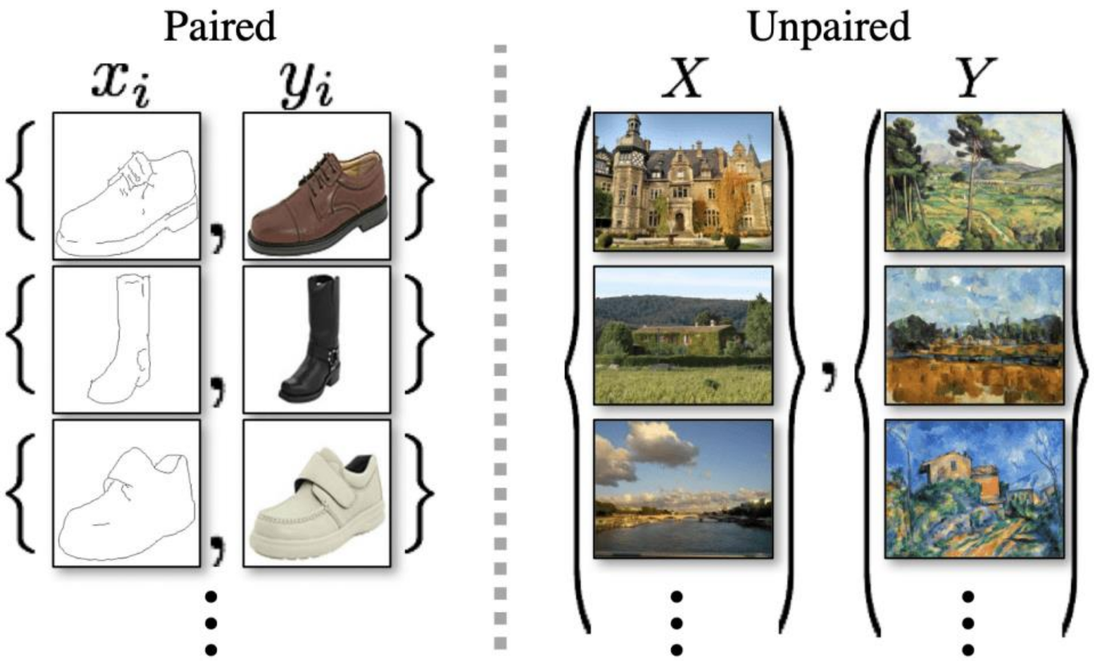

Image Source: https://pyimagesearch.com/2022/09/12/cyclegan-unpaired-image-to-image-translation-part-1/

Apply Generative Adversarial Networks

Necessity for Different Architectures: CycleGAN Introduction

- ✓ In unpaired image-to-image translation, conventional architectures such as GANs used in paired translations fall short due to the absence of direct mapping between image pairs.

- ✓ This is where **CycleGAN** comes in. CycleGAN was introduced to handle unpaired data using a cycle-consistency loss to ensure that the translation from one domain to another and back to the original domain results in an image similar to the original.

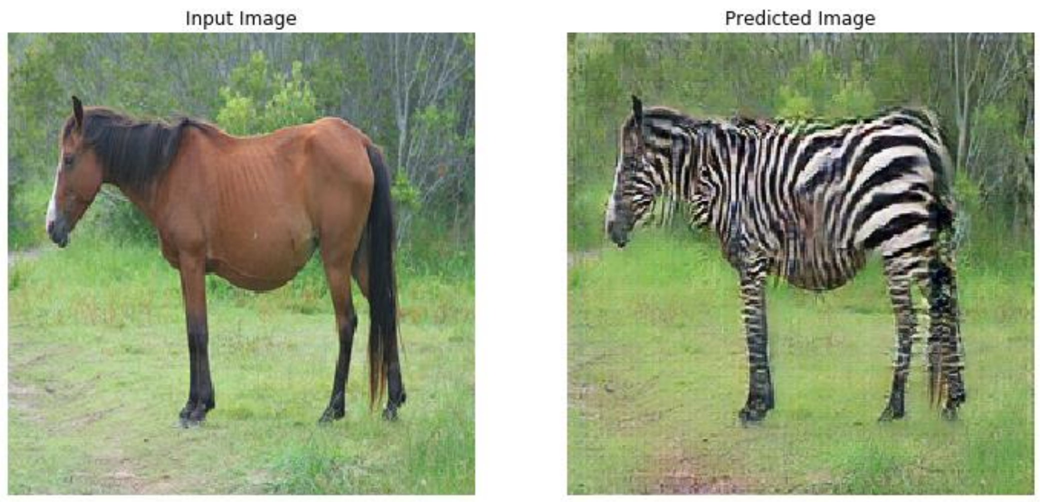

https://www.tensorflow.org/tutorials/generative/cyclegan

Apply Generative Adversarial Networks

CycleGAN Architecture:

- **Cycle Consistency Loss:** Ensures that if an image from domain A is translated to domain B, and then translated back to domain A, the output should be close to the original input image.
- **Generative Networks:** Two generators are used, one for each domain, to translate images from domain A to domain B and vice versa.
- **Discriminators:** Two discriminators are employed to ensure that the generated images are indistinguishable from the images in their respective domains.

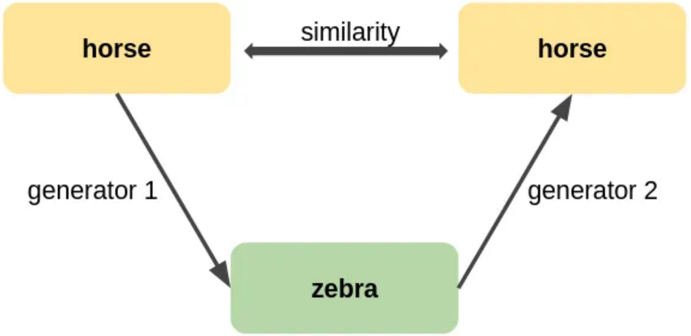

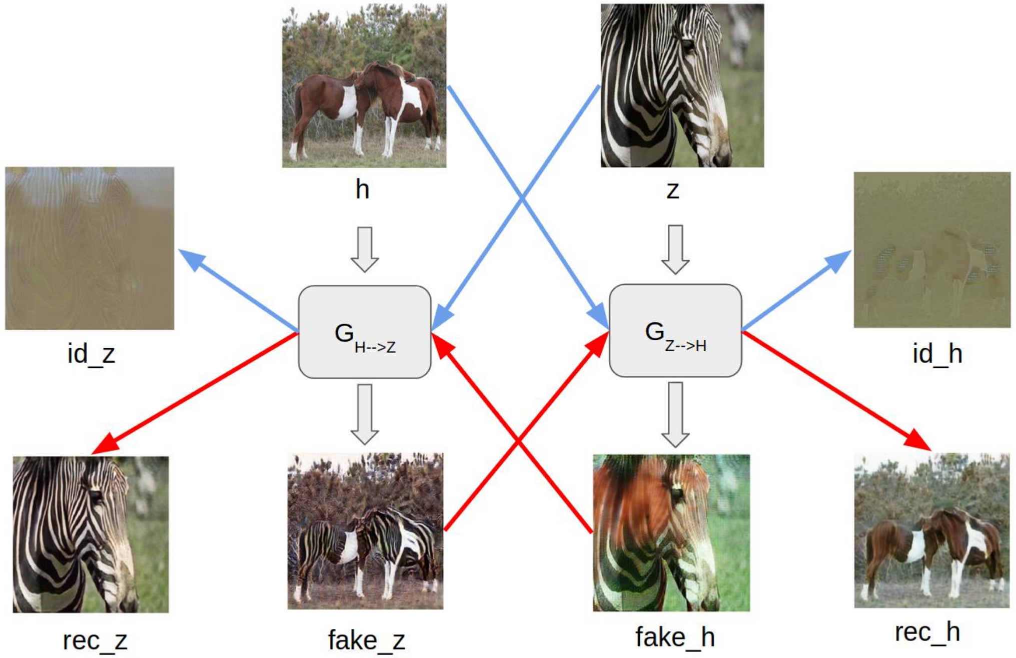

Image Source: https://towardsdatascience.com/overview-of-cyclegan-architecture-and-training-afee31612a2f

Apply Generative Adversarial Networks

CycleGAN Architecture:

Data : One training sample of the CycleGAN is formed by a random picture of horse and a random picture of a zebra.

Image Source :https://towardsdatascience.com/overview-of-cyclegan-architecture-and-training-afee31612a2f

Apply Generative Adversarial Networks

CycleGAN Architecture:

Model : Both generators have the following architecture (norm and activation layers aside) :

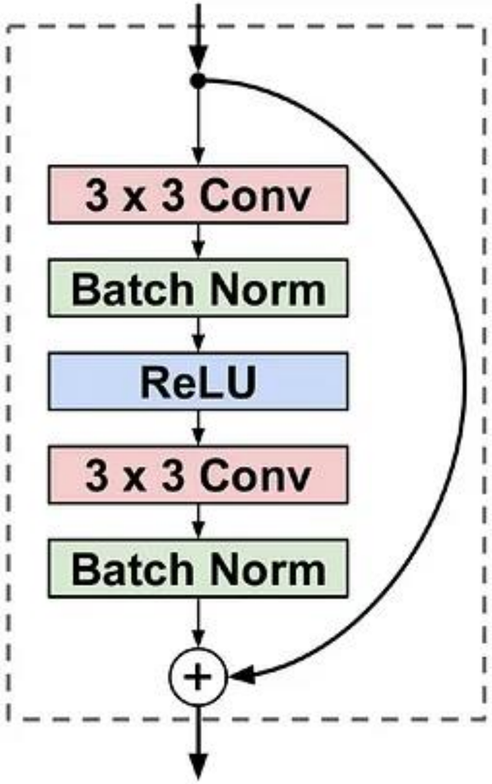

Layer	Activation Size
Input	3x256x256
64x7x7 conv, stride 1, pad 3	64x256x256
128x3x3 conv, stride 2, pad 1	128x128x128
256x3x3 conv, stride 2, pad 1	256x64x64
9 consecutive Residual Blocks	256x64x64
128x3x3 convTranspose, stride 2, pad 1, out_pad 1	128x128x128
64x3x3 convTranspose, stride 2, pad 1, out_pad 1	64x256x256
3x7x7 conv, stride 1, pad 3	3x256x256

Apply Generative Adversarial Networks

CycleGAN Architecture:

Result

✓ The CycleGAN is then trained for 200 epochs. During the first 100 epochs, the learning rate is fixed to 2e-4 and for the last 100 epochs the learning rate is linearly annealed from 2e-4 to 2e-6.

But there are also cases of failure (left: original images, right: generated images) :

Image Source :https://towardsdatascience.com/overview-of-cyclegan-architecture-and-training-afee31612a2f

Apply Generative Adversarial Networks

1. Dual Generative Adversarial Networks

In CycleGAN, there are **two GANs**, each consisting of:

Generator (G): The role of this network is to generate images that appear realistic from one domain to the other.

Discriminator (D): This network is responsible for distinguishing between real images from the domain and fake (generated) images.

The two GANs operate as follows:

GAN 1: Translates images from **domain X** to **domain Y**.
GAN 2: Translates images from **domain Y** to **domain X**.

Apply Generative Adversarial Networks

Dual Generative Adversarial Networks

Example:

Domain X: Images of horses.
Domain Y: Images of zebras.

- ✓ The first GAN learns to transform a horse image into a zebra.
- ✓ The second GAN learns to transform a zebra image into a horse.

Image Source https://www.analyticsvidhya.com/blog/2022/02/classification-without-training-data-zero-shot-learning-approach/

Apply Generative Adversarial Networks

Applications of Two GANs in One

- ✓ **Artistic Style Transfer:** Transforming photos into artistic renditions without paired examples.
- ✓ **Photo Enhancement:** Converting low-quality or stylized images into more realistic forms (e.g., black-and-white to color, day-to-night, etc.).
- ✓ **Domain Adaptation:** Translating between domains where data is inherently unpaired, such as converting satellite images to maps, transforming real-world photos into synthetic data for training models, or even medical image translations like MRI-to-CT scans.

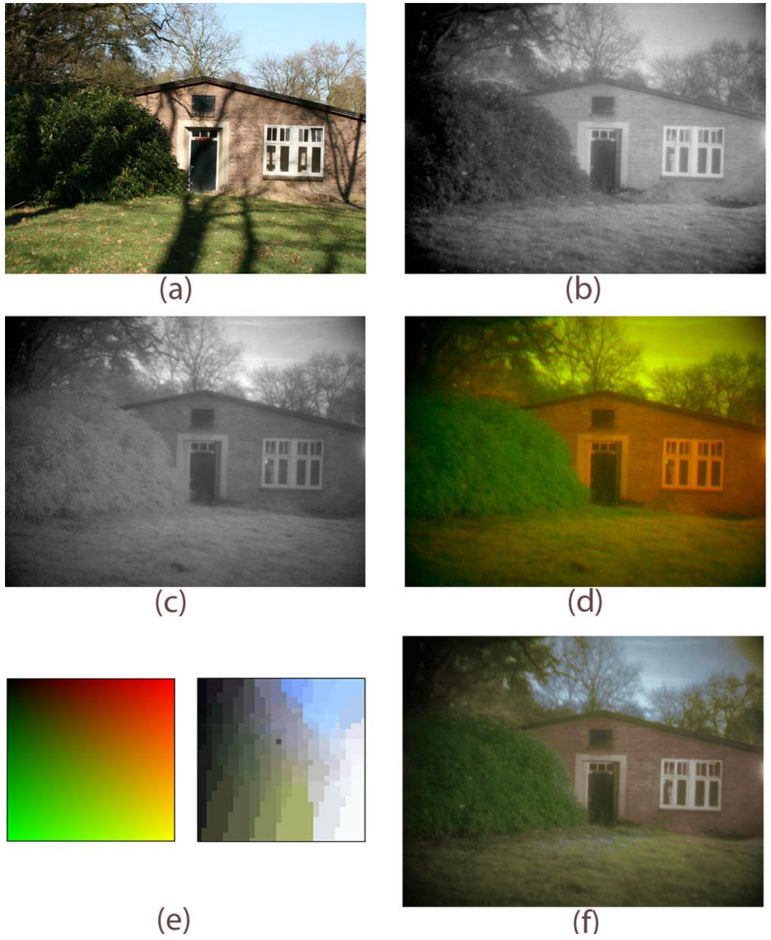

Image Source : https://spie.org/news/4447-color-remapping-turns-night-into-day#_=_

Comparison: Paired vs. Unpaired Image-to-Image Translation

Apply Generative Adversarial Networks

Comparison: Paired vs. Unpaired Image-to-Image Translation

Aspect	Paired Image Translation (Pix2Pix)	Unpaired Image Translation (CycleGAN)
Training Data	Requires aligned image pairs (source-target pairs)	No need for aligned image pairs; works with unpaired images from different domains
Architecture	Uses a single GAN model for translation from one domain to another	Uses two GANs with cycle-consistency (dual generators and discriminators)
Loss Functions	Relies on L1 loss and adversarial loss	Combines adversarial loss and cycle-consistency loss
Complexity	Simpler architecture since it only handles paired data	More complex due to dual GANs and the need to enforce cycle-consistency
Transformation Accuracy	Produces highly accurate transformations due to direct mapping	Slightly less accurate as there is no direct supervision (no paired data)
Training Stability	More stable due to paired supervision	More difficult to stabilize due to unpaired data and complex loss functions
Computation Requirements	Relatively lower computational cost	Higher computational cost due to the dual GAN setup and additional losses
Generalization Ability	Can overfit to specific paired data	Better generalization as it learns domain-level features

Apply Generative Adversarial Networks

Advantages and Disadvantages: Paired vs. Unpaired Image-to-Image Translation

Model	Advantages	Disadvantages
Pix2Pix (Paired)	Direct supervision via paired data improves accuracy	Requires large datasets of perfectly aligned image pairs
	Faster training due to clear mappings	Limited to domains where paired data is available
	Simpler architecture, easier to implement	Can be overfit to paired datasets, reducing generalization
CycleGAN (Unpaired)	No need for paired datasets	Harder to train due to unpaired data
	Better generalization to unseen data	More computationally expensive (two GANs, cycle-consistency)
	Suitable for many real-world tasks without aligned data	Slightly less accurate compared to paired models

Apply Generative Adversarial Networks

Use Cases for Each Model

Model	Use Cases
Pix2Pix (Paired)	- Image-to-image translation where paired data is available (e.g., satellite images to maps, photo enhancement)
	- **Medical Imaging:** translating between MRI and CT scans, where paired scans are provided
	- Semantic segmentation, where ground truth labels for images exist
CycleGAN (Unpaired)	- Artistic style transfer (e.g., photos to Van Gogh style without paired images)
	- Photo enhancements without ground truth (e.g., day-to-night, photo-to-sketch)
	- Translating between different domains with unpaired data (e.g., converting horses to zebras)

Case Study 3: Applications of GANs

Apply Generative Adversarial Networks

Case Study 3: Applications of GANs

1. Application Exploration with Respect to Data Augmentation, Privacy, and Anonymity

- **Data Augmentation:**
  - ✓ GANs are widely used to generate additional data samples from limited datasets.
  - ✓ They improve model training by artificially increasing dataset size, particularly in scenarios with imbalanced data.
  - ✓ **Example:** In medical imaging, GANs generate synthetic samples of rare diseases, enabling better diagnostic models without the need for vast real-world data.

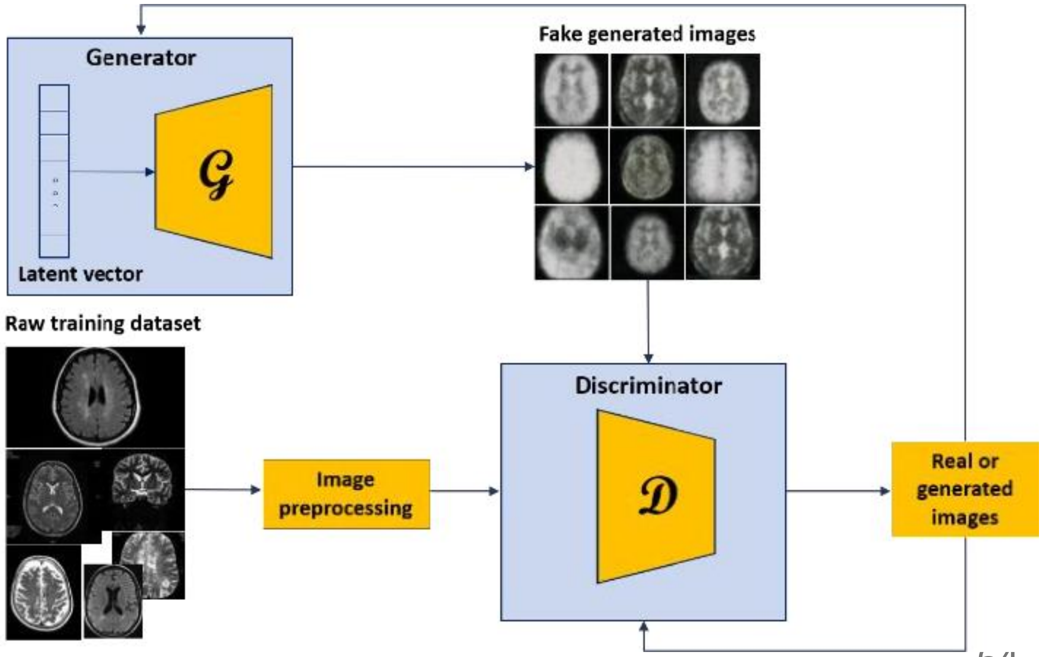

Image Source : https://www.researchgate.net/figure/Generative-adversarial-network-GAN-architecture-It-consists-of-two-sub-modules-a_fig1_368515565

Apply Generative Adversarial Networks

Case Study 3: Applications of GANs

- **Privacy Preservation:**
  - ✓ GANs help generate synthetic data that retains the statistical properties of real data without exposing sensitive information.
  - ✓ **Use case:** In healthcare, GANs create patient data surrogates, allowing researchers to work with valuable data without violating privacy regulations like HIPAA.

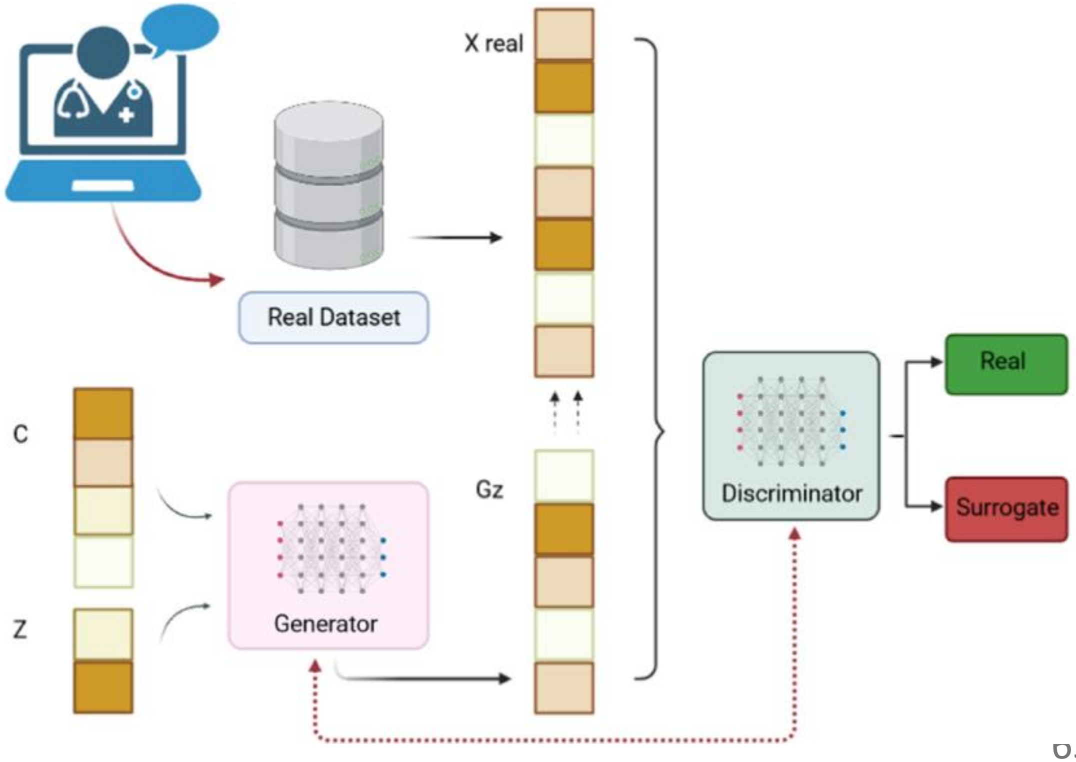

Image Source : https://www.researchgate.net/figure/n-a-conditional-generative-adversarial-network-cGAN-both-the-generator-G-and-the_fig1_371751970

Apply Generative Adversarial Networks

Case Study 3: Applications of GANs

- **Anonymity:**
  - ✓ GANs can help anonymize individuals in sensitive datasets while maintaining the utility of the data.
  - ✓ **Example:** Facial recognition datasets can anonymize features while keeping other attributes (e.g., expressions) intact, ensuring compliance with privacy laws while maintaining functionality.

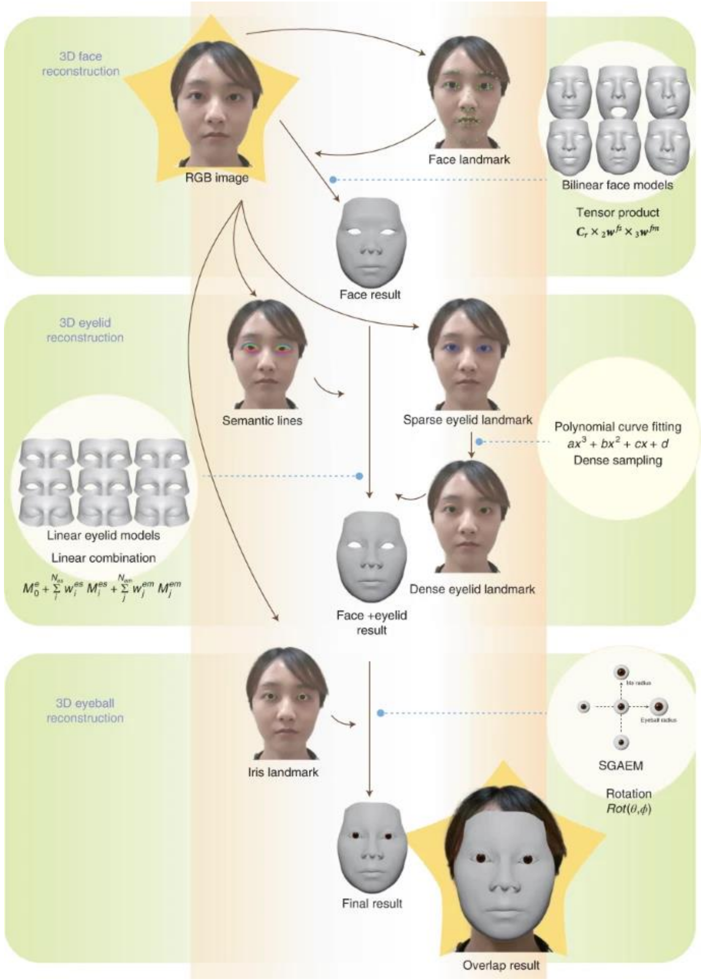

Image Source : https://www.nature.com/articles/s41591-022-01967-0

Apply Generative Adversarial Networks

Use Case Beyond Images: Application in Other Modalities

- **Text Data:**
  - ✓ GANs are employed to generate realistic text data, useful for natural language processing (NLP) applications, such as chatbot training, or creating synthetic datasets for language models.
  - ✓ **Example:** Text-to-text translation or text summarization using **SeqGAN** or **TextGAN**.
- **Audio Data:**
  - ✓ GANs can synthesize audio data, like generating realistic voice samples for speech synthesis or augmenting audio datasets for improved model performance.
  - ✓ **Example:** Voice conversion or speech enhancement tasks in fields like telecommunication or assistive technologies.
- **Video Data:**
  - ✓ GANs are used in generating and augmenting video data, such as creating realistic synthetic videos for surveillance, animation, or deepfake detection.
  - ✓ **Example:** Generating frame-by-frame synthetic data for tasks like human motion prediction.

Apply Generative Adversarial Networks

Implementations of Pix2Pix and CycleGAN

- **Pix2Pix Insights**
  - ✓ Works well for applications where paired data is available, such as transforming satellite images into map views or colorizing grayscale images.
  - ✓ Highly accurate transformations due to its reliance on paired data, making it suitable for tasks with clear mappings between input and output.

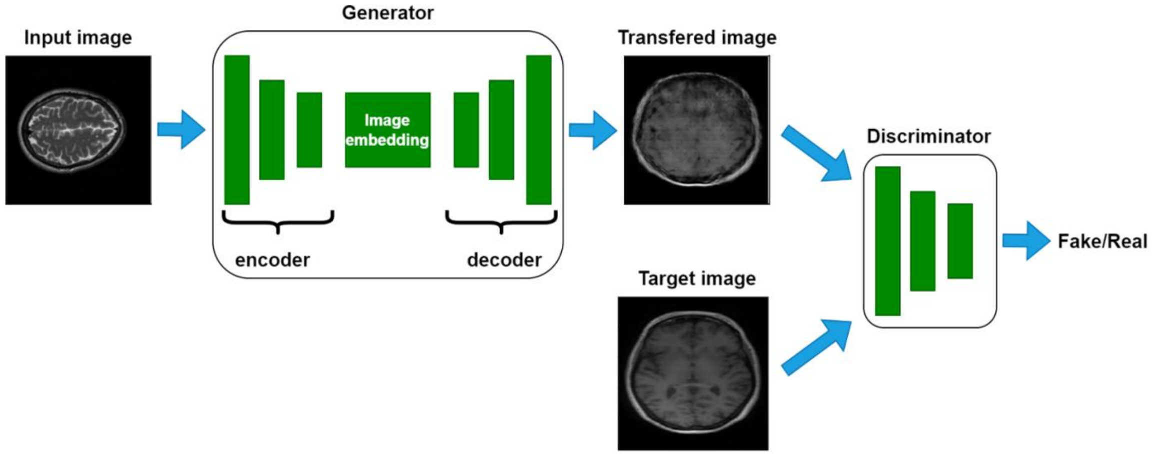

Apply Generative Adversarial Networks

Implementations of Pix2Pix and CycleGAN

- ✓ Ideal for scenarios where paired data is unavailable, such as style transfer (e.g., converting day photos to night photos or horses to zebras).
- ✓ Offers better generalization across domains without the need for direct mappings between images, enabling creative and practical applications in art, entertainment, and research.
- ✓ Its dual-GAN structure with cycle-consistency loss ensures that transformations can go back and forth, maintaining fidelity in the translation process.

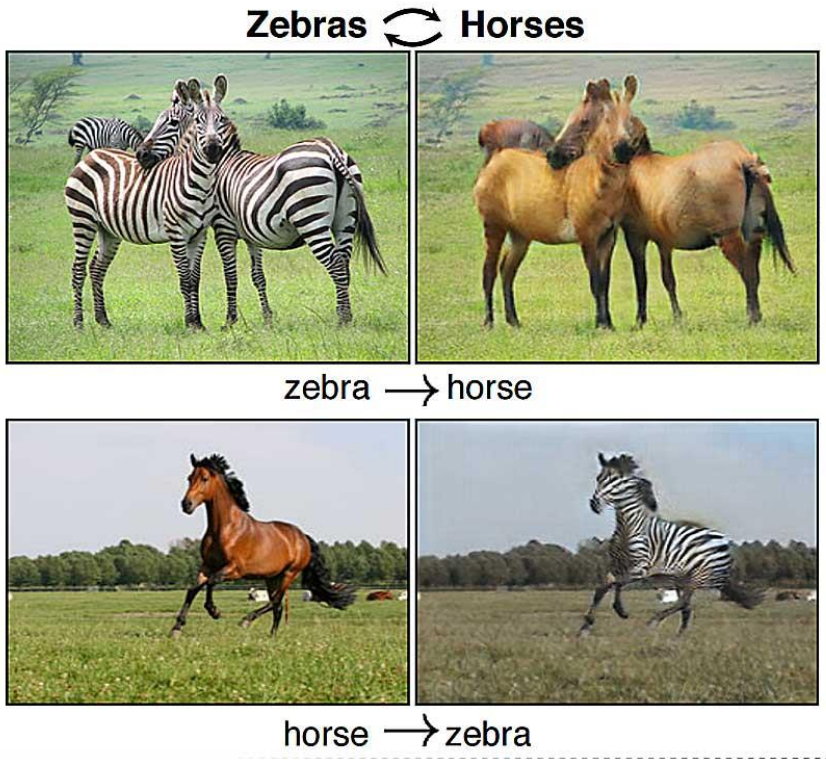

Image Source :https://medium.com/analytics-vidhya/gans-a-brief-introduction-to-generative-adversarial-networks-f06216c7200e

Recap of GAN Techniques

Apply Generative Adversarial Networks

Recap of GAN Techniques: Augmentation, Privacy, Anonymity

Data Augmentation with GANs

- ✓ GANs generate synthetic data that can be used to augment training datasets, improving model robustness and performance.
- ✓ **Applications:** Medical imaging (e.g., **rare disease synthesis**), facial recognition, and autonomous driving.

Privacy Preservation

- ✓ GANs create synthetic datasets that mimic real-world data without revealing sensitive or personal information.
- ✓ **Use case:** Protecting patient data in healthcare, or anonymizing individuals in large datasets without sacrificing utility.

Anonymity with GANs

- ✓ GANs can anonymize personally identifiable information (e.g., **faces, voices**) while preserving the underlying features needed for analysis.
- ✓ **Example:** Anonymizing faces in surveillance footage while maintaining expressions for emotion detection models.

Apply Generative Adversarial Networks

Recap of GAN Techniques

- **Key Takeaways: Image-to-Image Translation (Paired vs. Unpaired)**

- **Paired Image-to-Image Translation (Pix2Pix):**
  - ✓ Requires paired datasets (source and target images must be aligned).
  - ✓ Highly accurate due to the direct supervision, making it ideal for structured tasks like image segmentation, medical scans, and photo enhancement.
  - ✓ Simpler architecture and faster training but limited by the need for paired data.

Apply Generative Adversarial Networks

Recap of GAN Techniques

- **Unpaired Image-to-Image Translation (CycleGAN):**

  ✓ Works with unpaired datasets, eliminating the need for aligned image pairs.

  ✓ Effective for tasks where such pairing is impractical, like style transfer (day-to-night, horse-to-zebra, artistic transformations).

  ✓ Cycle-consistency loss ensures translation coherence, but the model is more complex and computationally expensive.

Apply Generative Adversarial Networks

Final Thoughts on the Future of GANs

- **Expanding Use Cases:**
  - ✓ GANs are not limited to image generation. They are being applied in diverse fields such as text synthesis, audio generation, video creation, and even 3D model generation.
- **Challenges and Innovations:**
  - ✓ While GANs offer incredible potential, challenges like training instability, mode collapse, and data bias remain. Ongoing research is focused on improving GAN architectures, including self-supervised learning and advanced loss functions.

Apply Generative Adversarial Networks

Final Thoughts on the Future of GANs..

- **Future Applications:**

  ✓ GANs hold great promise in industries like healthcare (drug discovery, medical image analysis), autonomous vehicles (simulation, synthetic data creation), and creative industries (art, music, animation).

  ✓ With advancements in privacy-preserving GANs and ethical AI, GANs will play a crucial role in responsible AI deployment, ensuring data privacy and unbiased model development.

Check your Knowledge

Apply Generative Adversarial Networks

Check your Knowledge

1. Which of the following is a primary application of GANs for data augmentation?

a) Reducing the size of datasets
b) Generating synthetic data to increase dataset diversity
c) Encrypting sensitive datasets
d) Removing noise from images

2. How do GANs help preserve privacy in sensitive data domains like healthcare?

a) By encrypting patient data
b) By generating synthetic data that mimics sensitive datasets
c) By directly replacing real data with dummy data
d) By anonymizing data through encryption

Apply Generative Adversarial Networks

Check your Knowledge

1. Which of the following is a primary application of GANs for data augmentation?

a) Reducing the size of datasets

b) Generating synthetic data to increase dataset diversity

c) Encrypting sensitive datasets

d) Removing noise from images

Answer: b) Generating synthetic data to increase dataset diversity

2. How do GANs help preserve privacy in sensitive data domains like healthcare?

a) By encrypting patient data

b) By generating synthetic data that mimics sensitive datasets

c) By directly replacing real data with dummy data

d) By anonymizing data through encryption

Answer: b) By generating synthetic data that mimics sensitive datasets

Apply Generative Adversarial Networks

Check your Knowledge

3. What is the main advantage of using GANs for anonymizing data?
a) They improve image resolution.
b) They retain relevant features while removing identifiable details.
c) They delete sensitive information permanently.
d) They use supervised learning techniques to anonymize data.

4. What does the image-to-image translation framework allow in GANs?
a) Transforming 2D images into 3D models
b) Translating input images from one domain into another domain
c) Enhancing image resolution
d) Applying filters to images for better visualization

Apply Generative Adversarial Networks

Check your Knowledge

3. What is the main advantage of using GANs for anonymizing data?
a) They improve image resolution.
b) They retain relevant features while removing identifiable details.
c) They delete sensitive information permanently.
d) They use supervised learning techniques to anonymize data.
Answer: b) They retain relevant features while removing identifiable details.

4. What does the image-to-image translation framework allow in GANs?
a) Transforming 2D images into 3D models
b) Translating input images from one domain into another domain
c) Enhancing image resolution
d) Applying filters to images for better visualization
Answer: b) Translating input images from one domain into another domain

Apply Generative Adversarial Networks

Check your Knowledge

5. What is the primary architecture of Pix2Pix used for paired image-to-image translation?
a) LSTM and CNN
b) Transformer and CNN
c) U-Net Generator and PatchGAN Discriminator
d) Fully connected neural network

6. Why is paired data essential for the Pix2Pix model?
a) It ensures unsupervised learning.
b) It provides a one-to-one correspondence between input and output images.
c) It removes noise from the data.
d) It is required for generating high-resolution images.

Apply Generative Adversarial Networks

Check your Knowledge

5. What is the primary architecture of Pix2Pix used for paired image-to-image translation?
a) LSTM and CNN
b) Transformer and CNN
c) U-Net Generator and PatchGAN Discriminator
d) Fully connected neural network
Answer: c) U-Net Generator and PatchGAN Discriminator

6. Why is paired data essential for the Pix2Pix model?
a) It ensures unsupervised learning.
b) It provides a one-to-one correspondence between input and output images.
c) It removes noise from the data.
d) It is required for generating high-resolution images.
Answer: b) It provides a one-to-one correspondence between input and output images.

Apply Generative Adversarial Networks

Check your Knowledge

7. What is the main difference between paired and unpaired image-to-image translation?

a) Paired translation uses two generators, while unpaired translation uses one.
b) Paired translation requires a dataset where input images have corresponding outputs, whereas unpaired translation does not.
c) Unpaired translation only works with supervised learning, while paired translation is unsupervised.
d) Unpaired translation works only on images, while paired translation works on videos.

8. In CycleGAN, what ensures the output image can be transformed back into the original input?

a) Convolutional layers
b) Adversarial loss
c) Cycle consistency loss
d) LSTM layers

Apply Generative Adversarial Networks

Check your Knowledge

7. What is the main difference between paired and unpaired image-to-image translation?

a) Paired translation uses two generators, while unpaired translation uses one.
b) Paired translation requires a dataset where input images have corresponding outputs, whereas unpaired translation does not.
c) Unpaired translation only works with supervised learning, while paired translation is unsupervised.
d) Unpaired translation works only on images, while paired translation works on videos.

Answer: b) Paired translation requires a dataset where input images have corresponding outputs, whereas unpaired translation does not.

8. In CycleGAN, what ensures the output image can be transformed back into the original input?

a) Convolutional layers
b) Adversarial loss
c) Cycle consistency loss
d) LSTM layers

Answer: c) Cycle consistency loss

Apply Generative Adversarial Networks

Check your Knowledge

9. Which of the following is a use case of unpaired image-to-image translation with CycleGAN?

a) Transforming satellite images into map routes
b) Translating horses to zebras and vice versa
c) Converting handwritten text to typed text
d) Colorizing black-and-white images

10. What are the key components of CycleGAN architecture?

a) Two generators and two discriminators
b) One generator and one discriminator
c) U-Net generator and PatchGAN discriminator
d) Three generators and two discriminators

Apply Generative Adversarial Networks

Check your Knowledge

9. Which of the following is a use case of unpaired image-to-image translation with CycleGAN?

a) Transforming satellite images into map routes

b) Translating horses to zebras and vice versa

c) Converting handwritten text to typed text

d) Colorizing black-and-white images

Answer: b) Translating horses to zebras and vice versa

10. What are the key components of CycleGAN architecture?

a) Two generators and two discriminators

b) One generator and one discriminator

c) U-Net generator and PatchGAN discriminator

d) Three generators and two discriminators

Answer: a) Two generators and two discriminators

?

Thank You !!!